

Terraform State

Module Topics

- ◆ Understanding terraform state
- ◆ Managing state with terraform state commands
- ◆ Local state backend risks
- ◆ Using workspaces to multiple configurations
- ◆ Remote backends - using AWS S3
- ◆ Migrating state backends

What Is Terraform State?

- ◆ Terraform state refers to the record terraform keeps of the resources it has created
- ◆ The state file:
 - Creates a mapping between what was specified in the *.tf files and deployed AWS resources
 - Stores metadata and configuration information
- ◆ When *terraform plan* is run, it does the following
 - Reads all of the *.tf files in a directory
 - Updates the state information to record modifications from the *.tf files
 - Queries AWS to get a description the current state of the deployed resources
 - Creates a plan to modify the AWS resources to conform to the state descriptions
- ◆ Terraform cannot see or modify AWS resources that are not in its state file

Terraform Planning

- ◆ When *terraform plan* is run
 - All files in the current directory with the .tf extension are read and their contents merged
 - Terraform develops a plan for implementing the specified configuration
- ◆ A directed acyclic graph of operations is created to ensure that all dependencies are resolved
 - The resources may need to be created in a specific order
- ◆ Thus means we can sort organize our code anyway we want
- ◆ However, there is usually a standard way to do this - i.e. the canonical form introduced in the last module

Terraform is Declarative

- ◆ Declarative means that you only describe the final state that you want your AWS resources to be in
- ◆ After you run *terraform apply* you will see the output in *terraform.tfstate*

```
{
  "version": 4,
  "terraform_version": "0.12.0",
  "serial": 1,
  "lineage": "1f2087f9-4b3c-1b66-65db-8b78faafc6fb",
  "outputs": {},
  "resources": [
    {
      "mode": "managed",
      "type": "aws_instance",
      "name": "example",
      "provider": "provider.aws",
      "instances": [
        {
          "schema_version": 1,
          "attributes": {
            "ami": "ami-0c55b159cbfafa1f0",
            "availability_zone": "us-east-2c",
            "id": "i-00d689a0acc43af0f",
            "instance_state": "running",
            "instance_type": "t2.micro",
            "(...)": "(truncated)"
          }
        }
      ]
    }
  ]
}
```

Meaning of the "terraform.tfstate" on Previous Slide

- ◆ Resource with *type aws_instance* and *name example* corresponds to an EC2 Instance in the AWS account with ID i-00d689a0acc43af0f
- ◆ Every time you run Terraform
 - it can fetch the latest status of this EC2 Instance from AWS
 - compare that to what's in your Terraform configurations
 - determine what changes need to be applied
- ◆ Thus, the output of the *terraform plan* command is a diff
 - between the code on your computer and
 - the infrastructure deployed in the real world, as discovered via IDs in the state file

Example - Understanding State

- ◆ In this example, we have defined two EC2 instances "x" and "y" in our main.tf file
 - Since terraform plan hasn't been run, there is no terraform state
 - Note there is already a running EC2 instance in our AWS account


Local *.tf Files

```
resource "aws_instance" "x" {  
  type = "t2.micro"}  
  
resource "aws_instance" "y" {  
  type = "t2.micro"}
```

Terraform State



Running AWS Resources



Before "terraform plan"

Example - Understanding State - The code

- ◆ This screenshot shows the code being used for this example

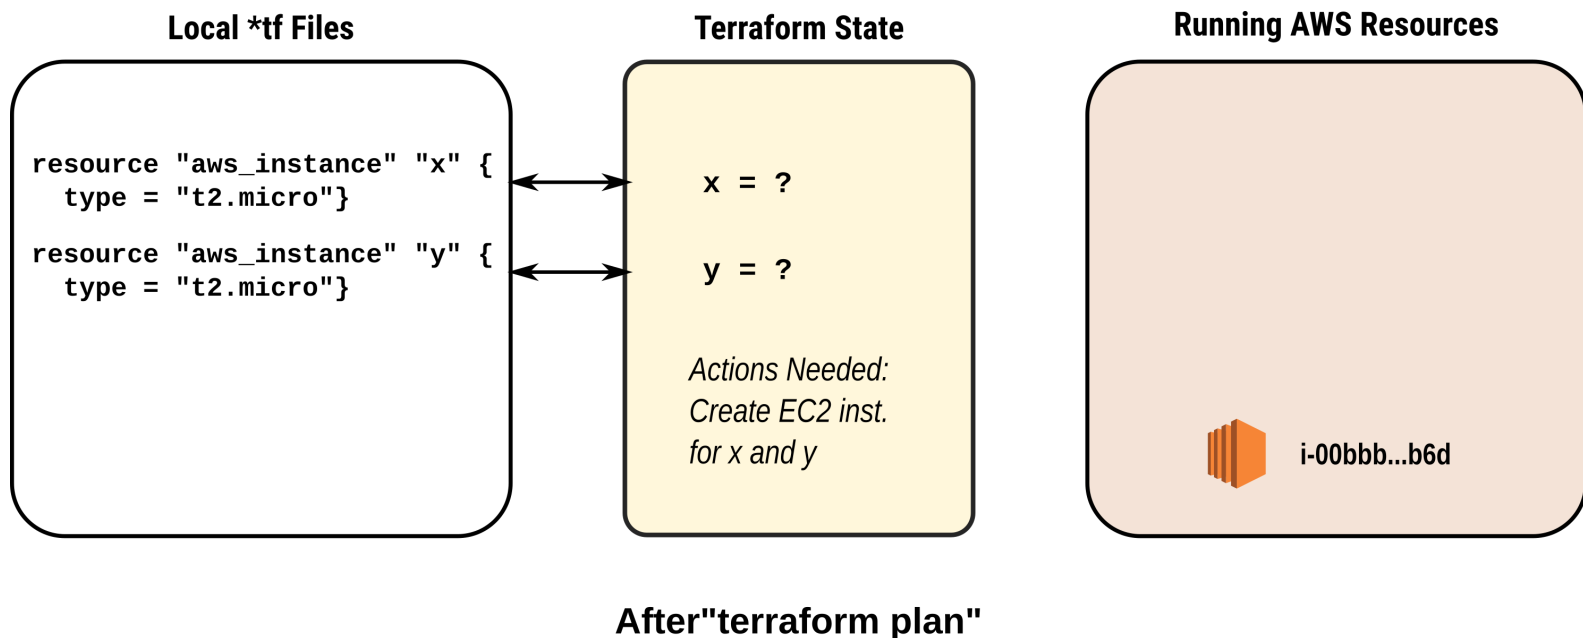
```
# Example 03-01

resource "aws_instance" "X" {
  ami = "ami-077e31c4939f6a2f3"
  instance_type = "t2.micro"
  tags = {
    Name = "Instance X"
  }
}

resource "aws_instance" "Y" {
  ami = "ami-077e31c4939f6a2f3"
  instance_type = "t2.micro"
  tags = {
    Name = "Instance Y"
  }
}
```

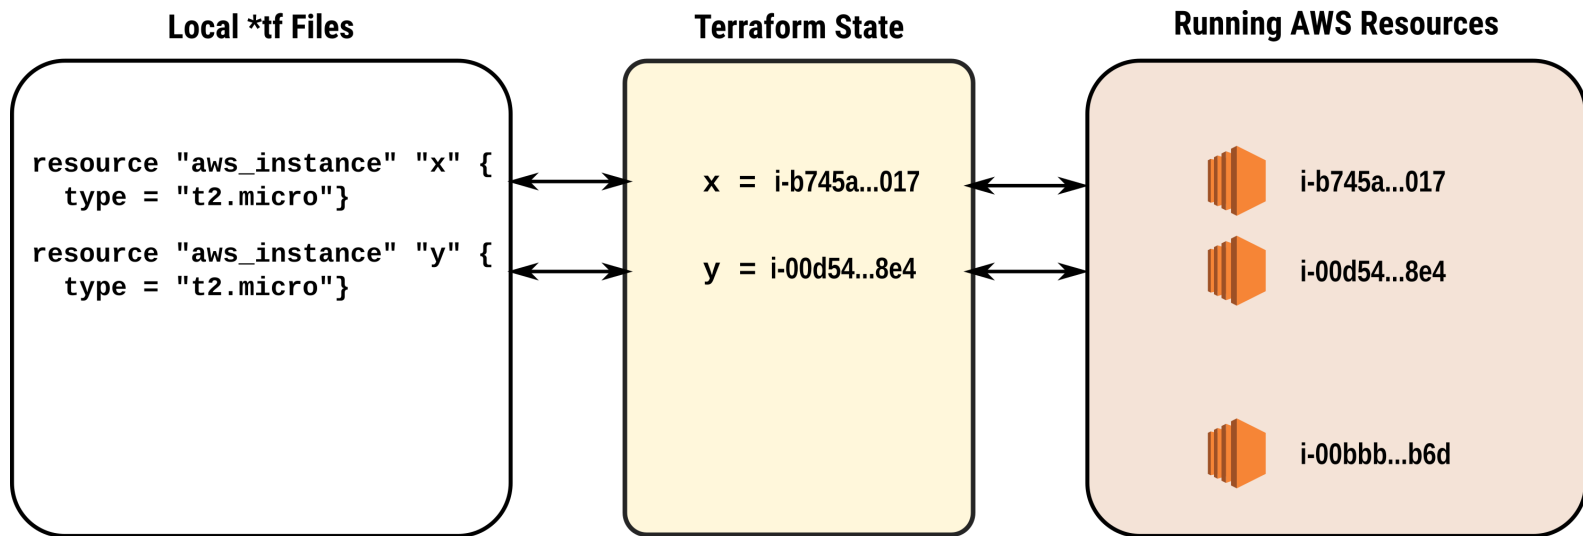

Example: The Plan Operation

- ◆ After running *terraform plan* there are two resources without corresponding AWS instances
 - Terraform writes an action plan that will bring the AWS environment into alignment with the terraform state
 - If actions must be done in a particular order, terraform will determine the correct sequence



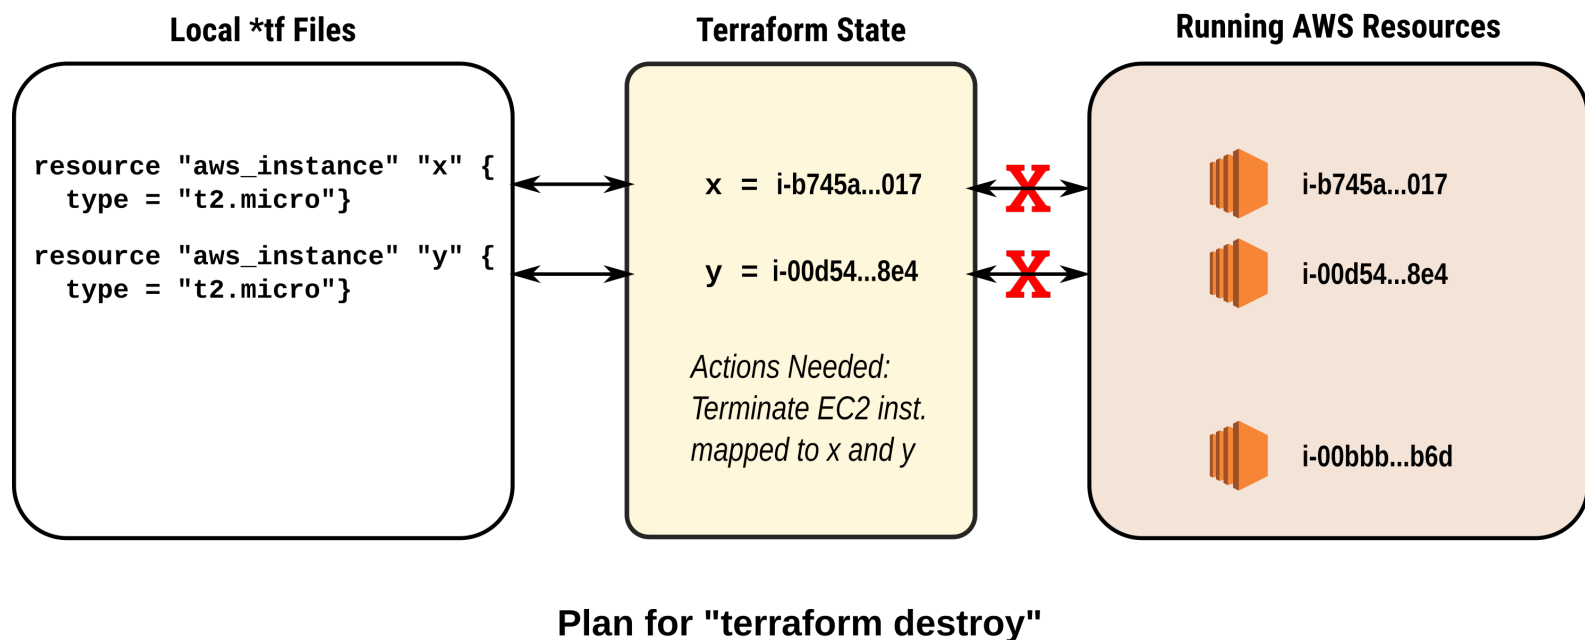
Example: After Apply Operation

- ◆ The required modifications are made to AWS and the results stored in the state file



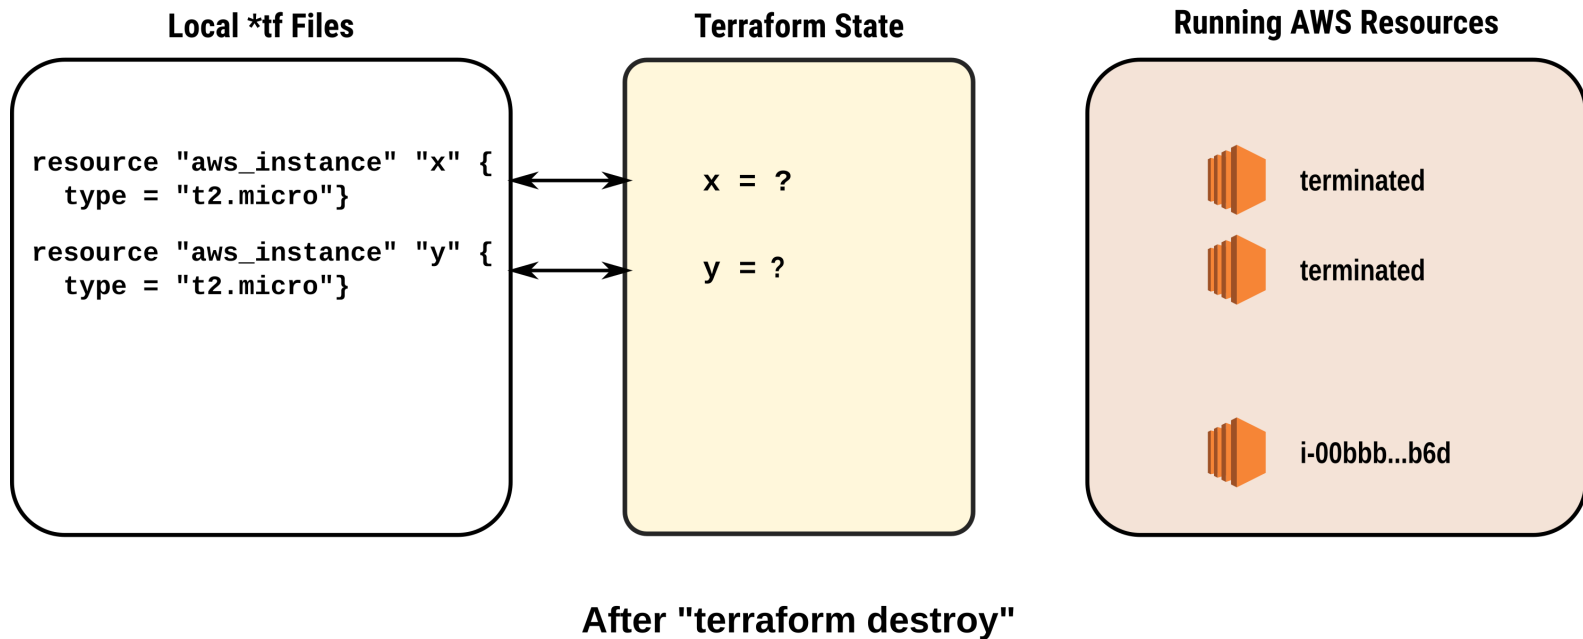
Example: Planning the Destroy Operation

- ◆ The only resources that will be destroyed are the ones that are actually in the terraform state file
 - Terraform plans the actions to terminate the resources it is managing
 - A with apply, terraform will determine the correct sequence for removing resources



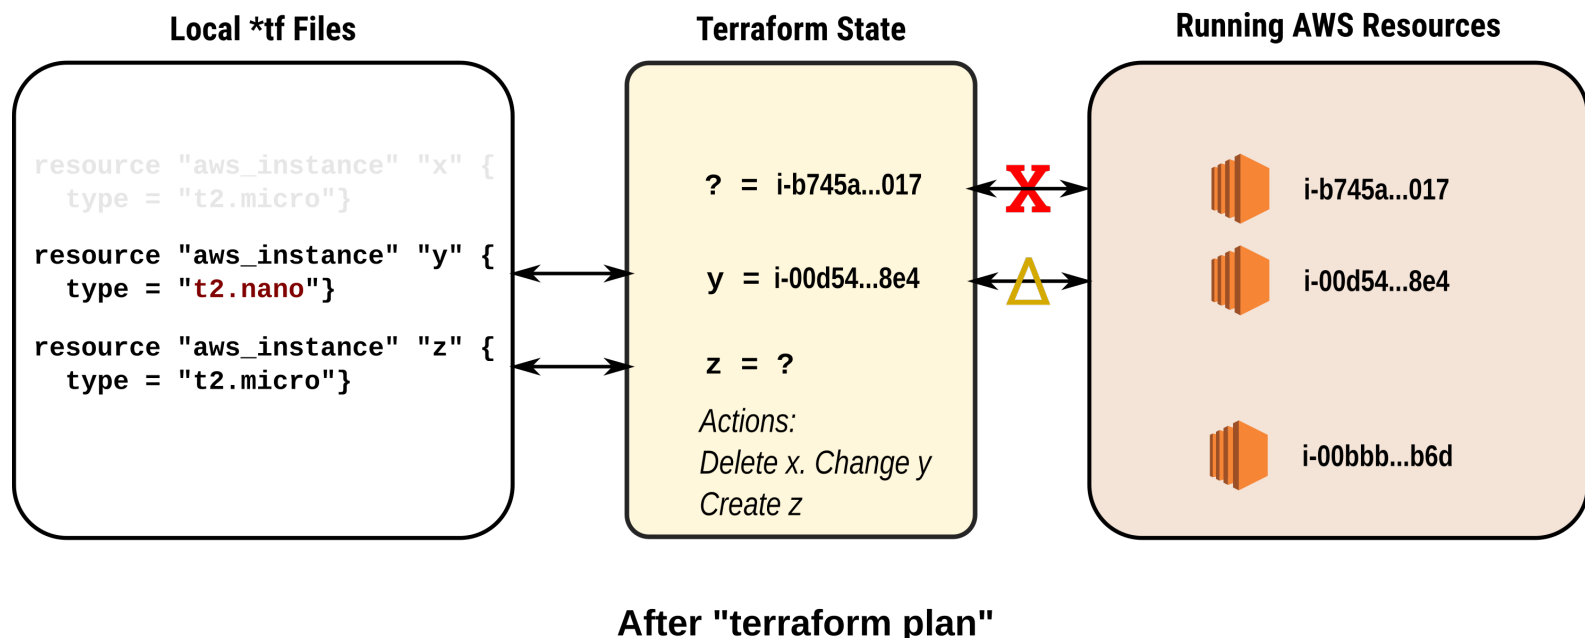
Example: After the Destroy Operation

- Any resources not listed in the terraform state are left untouched



Modifying a Resource

- ◆ Changes to the *.tf files are translated into actions by *terraform apply*
 - Removing a resource from the file causes its deletion from AWS and the state file
 - Adding a resource to the file causes it be created
 - Changing parameters of a resource causes it to be modified
 - **If a resource cannot be modified (eg. changing the ami) then the existing resource is destroyed and new resource created**



Lab 3-1

- ◆ Please do Lab 3-1

The "state" Command

- ◆ The state command has multiple options (not all are listed)
 - *terraform state list* : lists the resources being managed
 - *terraform state show <resource>* : displays state data for a resource
 - *terraform state rm <resource>* : stops managing the AWS object linked to <resource>
 - *terraform state mv* : renames a resource in the state file
- ◆ *terraform import <resource> <AWS ID>*: links the Terraform resource with a terraform resource

State Command Example

- ◆ The following is the environment defined for the example:

```
# Example 03-01

resource "aws_instance" "X" {
  ami = "ami-077e31c4939f6a2f3"
  instance_type = "t2.micro"
  tags = {
    Name = "Instance X"
  }
}

resource "aws_instance" "Y" {
  ami = "ami-077e31c4939f6a2f3"
  instance_type = "t2.micro"
  tags = {
    Name = "Instance Y"
  }
}
```


The "state list" Command

- ◆ The *state list* command lists all the resources being managed by the state file

```
C:\home\terraform>terraform state list  
aws_instance.X  
aws_instance.Y
```

The "state show" command

- ◆ The *state show* <id> displays the state file JSon data for that resource

```
C:\home\terraform>terraform state show aws_instance.X
# aws_instance.X:
resource "aws_instance" "X" {
    ami                        = "ami-077e31c4939f6a2f3"
    arn                      = "arn:aws:ec2:us-east-2:983803453537:ins
    associate_public_ip_address = true
    availability_zone        = "us-east-2c"
    cpu_core_count           = 1
    cpu_threads_per_core     = 1
    disable_api_termination  = false
    ebs_optimized            = false
    get_password_data        = false
    hibernation              = false
    id                      = "i-087c4d43d292c1c1b"
    instance_initiated_shutdown_behavior = "stop"
```

The "state rm" Command

- ◆ The rm command removes a specific resource from the state file
- ◆ In our example, we can remove the instance "X" from the state file
 - This means that the AWS resource is now no longer managed by terraform

```
C:\home\terraform>terraform state list
aws_instance.X
aws_instance.Y

C:\home\terraform>terraform state rm aws_instance.X
Removed aws_instance.X
Successfully removed 1 resource instance(s).

C:\home\terraform>terraform state list
aws_instance.Y
```

- ◆ If *terraform apply* is run again, a new version of aws_instance.X will be created because terraform can no longer 'see' the AWS instance it previously created

The "terraform import" Command

- ◆ This is the converse of the "state rm" command by moving an existing AWS resource into a state file
- ◆ There must be a terraform resource specification with parameters that match the properties of the existing AWS resource
- ◆ In this example, we add back aws_instance.X that we just removed

```
C:\home\terraform>terraform state list
aws_instance.Y

C:\home\terraform>terraform import aws_instance.X i-087c4d43d292c1c1b
aws_instance.X: Importing from ID "i-087c4d43d292c1c1b"...
aws_instance.X: Import prepared!
  Prepared aws_instance for import
aws_instance.X: Refreshing state... [id=i-087c4d43d292c1c1b]

Import successful!

The resources that were imported are shown above. These resources are now in
your Terraform state and will henceforth be managed by Terraform.

C:\home\terraform>terraform state list
aws_instance.X
aws_instance.Y
```

The "state mv" Command

- ◆ This command allows us to link an existing AWS resource to a different terraform specification
 - For example, if we want to rename our EC2 instance from aws_instance.X to aws_instance.Z

```
resource "aws_instance" "X" {
  ami = "ami-077e31c4939f6a2f3"
  instance_type = "t2.micro"
  tags = {
    Name = "Instance X"
  }
}

resource "aws_instance" "Z" {
  ami = "ami-077e31c4939f6a2f3"
  instance_type = "t2.micro"
  tags = {
    Name = "Instance Z"
  }
}
```

The "state mv" Command

- ◆ Executing the " terraform state mv" command breaks the association between aws_instance.X and the AWS resource and then re-associates it with aws_instance.Z

```
C:\home\terraform>terraform state list
aws_instance.X
aws_instance.Y

C:\home\terraform>terraform state mv aws_instance.X aws_instance.Z
Move "aws_instance.X" to "aws_instance.Z"
Successfully moved 1 object(s).

C:\home\terraform>terraform state list
aws_instance.Y
aws_instance.Z
```

The "state mv" Command

- ◆ Running the "terraform plan" shows that two actions now have to be taken to update the AWS environment
 - The new instance "Z" has to be updated to change the tag from "Resource X" to "Resource Z"
 - Since there is no longer an AWS instance associated with `aws_instance.Z`, a new AWS instance will have to be created

```
C:\home\terraform>terraform plan
aws_instance.Y: Refreshing state... [id=i-09dd6f80869afb6cf]
aws_instance.Z: Refreshing state... [id=i-087c4d43d292c1c1b]

Terraform used the selected providers to generate the following execution plan.
Resource actions are indicated with the
following symbols:
+ create
~ update in-place

Terraform will perform the following actions:

# aws_instance.X will be created
+ resource "aws_instance" "X" {
  + ami              = "ami-077e31c4939f6a2f3"
  + availability_zone = (known after apply)
```

```
# aws_instance.Z will be updated in-place
~ resource "aws_instance" "Z" {
  id              = "i-087c4d43d292c1c1b"
  ~ tags          = {
    ~ "Name" = "Instance X" -> "Instance Z"
  }
  ~ tags_all      = {
    ~ "Name" = "Instance X" -> "Instance Z"
  }
}
```

Tainting and Untainting

- ◆ Occasionally, an AWS resource is created but is degraded or damaged, often because of a transient AWS problem
 - Although the resource is created, it is in a suspicious state and is marked by terraform as being tainted
 - A tainted resource will be recreated the next time *terraform apply* is run
- ◆ You can also manually taint a resource by running the *terraform taint* command if you feel the resource should be recreated
- ◆ Any tainted resource can be untainted by running the *terraform untaint* command

Tainting and Untainting

```
C:\home\terraform>terraform taint aws_instance.X
Resource instance aws_instance.X has been marked as tainted.

C:\home\terraform>terraform plan
aws_instance.X: Refreshing state... [id=i-0f16d218c41248b7e]
aws_instance.Y: Refreshing state... [id=i-09dd6f80869afb6cf]

Terraform used the selected providers to generate the following execution
plan. Resource actions are indicated with the following symbols:
-/+ destroy and then create replacement

Terraform will perform the following actions:

# aws_instance.X is tainted, so must be replaced
-/+ resource "aws_instance" "X" {
```

```
C:\home\terraform>terraform untaint aws_instance.X
Resource instance aws_instance.X has been successfully untainted.

C:\home\terraform>terraform plan
aws_instance.Y: Refreshing state... [id=i-09dd6f80869afb6cf]
aws_instance.X: Refreshing state... [id=i-0f16d218c41248b7e]

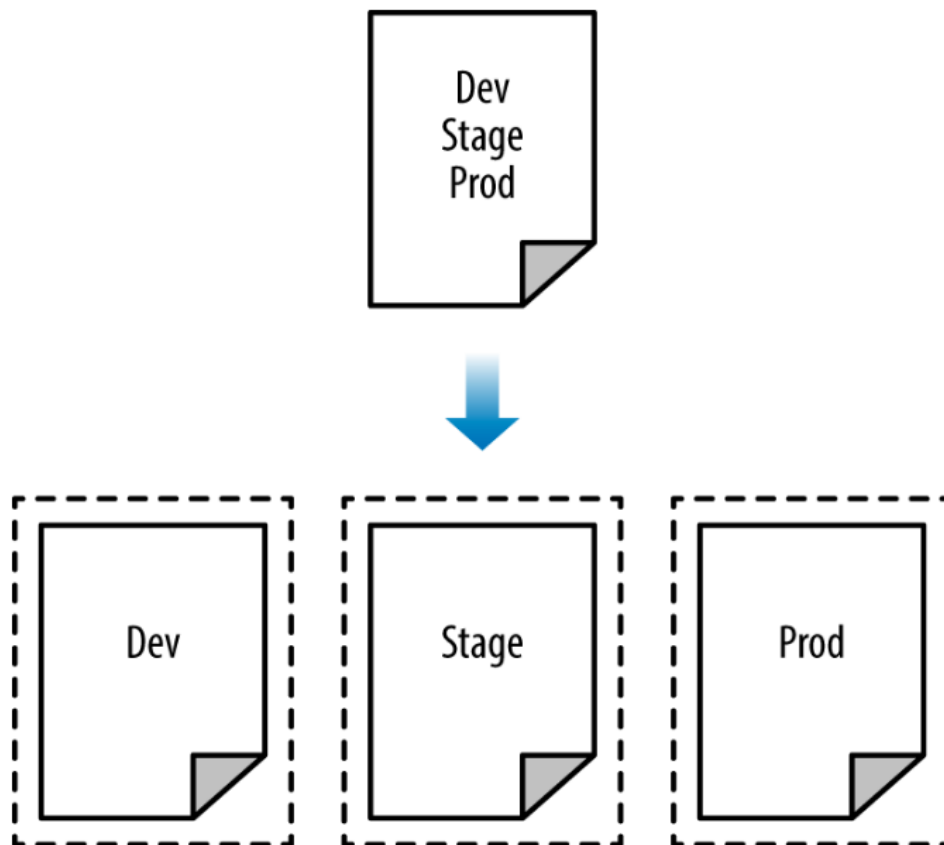
No changes. Infrastructure is up-to-date.
```

Lab 3-2

- ◆ Please do lab 3-2


Separate Environments

- ◆ We often need multiple copies of a deployment for different purposes
- ◆ Common environments are: development, test, stage and production



Terraform Workspaces

- ◆ Terraform supports a separate configuration for each deployment
 - Each deployment is called a workspace
 - There is always a *default* workspace
- ◆ We can create additional workspaces as we need them
 - For example, we could have defined dev and test workspaces



The "workspace" Command

- ◆ *terraform workspace* has several options:
 - *list* : lists all workspaces marking the current one with "*"
 - *show* : lists the currently active workspace
 - *new <name>* : creates and switches to the newly created workspace
 - *select <name>* : switches to the named workspace
 - *delete <name>* : deletes the named workspace
 - The "default" workspace can never be deleted
 - Deleting a workspace does **not** destroy the resources, it just leaves them unmanaged

Managing State versus Managing IaC Code

- ◆ Workspaces manage the state information of different but related deployments
 - Often useful for a quick "what-if" type of analysis
- ◆ However, state information is **not** IaC code
- ◆ In order to support separate environments, there are two issues that must be managed
 - The states of the deployments - terraform does that for us
 - The organization of the IaC source code in the *.tf files
- ◆ We have to manage the terraform source files exactly like we manage source code in standard development projects
 - For example, we can use a git-github model and DevOps continuous integration

Local Backends and Workspaces

- ◆ The location of the terraform state files is called the "backend"
- ◆ When the state file is kept in the same directory as the *.tf files, then we are using what is called a local backend
 - This is the default for terraform "out of the box"
- ◆ Each workspace manages its own copy of the AWS resources, but they all use the same *.tf files
 - The amount of "isolation" between our different workspaces or teams is quite low
 - For example, the `dev` group might make changes that break the `prod` configuration
- ◆ We often use workspaces when we want to spin up a copy of an environment without interfering with an existing deployment

Versioning Configurations

- ◆ A basic principle of IaC is that we treat our configuration files as code
- ◆ We can version our configuration by putting the *.tf files in git or other vcs
 - We do **not** put the state files in version control
 - Unless we want to store snapshots of the state files
- ◆ If we make changes and break a configuration, we can roll back to a working version
- ◆ If we want make changes in workspaces:
 - Commit the *.tf files to a git repository
 - For each workspace used, create a corresponding git branch
 - As you switch workspaces, switch branches
- ◆ This still does not fix the problems with the state files
 - *"Oops, I forgot to check out the dev branch and destroyed the production configuration"*

Problem with Local Backends

- ◆ Shared storage for state files
 - Files need to be in common shared area so everyone on the team can access them
 - Without file locking, race conditions when concurrent updates to the state files take place
 - This can lead to conflicts, data loss, and state file corruption
 - Even if we use versioning, branches and workspaces
- ◆ Isolation
 - It's difficult to isolate the code used in different environments
 - Lack of isolation makes it easy to accidentally overwrite environments
 - The problem that we cannot address locally is that the state file is a shared resource
 - Even if the *.tf source files are isolated from each other
- ◆ Secrets
 - Confidential information is stored in the clear (i.e. AWS Keys)

Remote Backends

- ◆ Each Terraform configuration has a location where the state files are kept
 - This is called the "backend"
 - The default is to use files in the local directory
 - Even using git or another system does not address the problems mentioned
- ◆ Terraform can also support "remote" backends
 - For example, we can keep state files in an S3 bucket on AWS
 - Not all providers can host remote back ends

Remote AWS Backend

- ◆ Using S3 as a backend resolves many of these issues
 - S3 manages the updating and access independently, and supports versions
 - S3 supports encryption
 - S3 supports locking for multiple access
 - S3 allows a common repository we can control access to
- ◆ S3 is also managed so that we don't have to manage it
 - S3 has high levels of availability and durability
 - S3 also means we have reduced the risk of "loosing" configurations

Setting up the S3 Bucket

```
resource "aws_s3_bucket" "terraform_state" {
  bucket = "terraform-up-and-running-state"

  # Prevent accidental deletion of this S3 bucket
  lifecycle {
    prevent_destroy = true
  }

  # Enable versioning so we can see the full revision history of our
  # state files
  versioning {
    enabled = true
  }

  # Enable server-side encryption by default
  server_side_encryption_configuration {
    rule {
      apply_server_side_encryption_by_default {
        sse_algorithm = "AES256"
      }
    }
  }
}
```

Setting up the Locking Table

- ◆ Next, a DynamoDB table to use for locking
 - DynamoDB is Amazon's distributed key-value store
 - It supports strongly consistent reads and conditional writes

```
resource "aws_dynamodb_table" "terraform_locks" {  
  name           = "terraform-up-and-running-locks"  
  billing_mode   = "PAY_PER_REQUEST"  
  hash_key       = "LockID"  
  
  attribute {  
    name = "LockID"  
    type = "S"  
  }  
}
```

The Backend for the Remote Backend

- ◆ When we set up the remote backend, we create a state file that describes the configuration of the remote backend
- ◆ **The remote backend state file is not kept in the remote backend**
 - We keep the remote backend state file separate and secure
 - Locked down and accessible only to the configuration manager
- ◆ We can have multiple S3 back ends for different projects
 - The state files for each S3 backend are kept in a master S3 backend
 - But, the state of the master S3 backend is stored securely somewhere else

Setting Up the Backend

- ◆ We have to tell Terraform the backend is now remote
 - We do this in the *terraform* directive
 - The key creates a unique folder in the S3 bucket for this state file

```
terraform {  
  backend "s3" {  
    # Replace this with your bucket name!  
    bucket      = "terraform-up-and-running-state"  
    key         = "mykey" #  
    region     = "us-east-2"  
  
    # Replace this with your DynamoDB table name!  
    dynamodb_table = "terraform-up-and-running-locks"  
    encrypt        = true  
  }  
}
```

Moving State File Locations

- ◆ To move local state to a remote backend
 - Create the remote backend resources and define the backend configuration
 - Run *terraform init* and the local config is copied to the remote backend
- ◆ To move from remote backend to a local backend
 - Remove the backend configuration
 - Run *terraform init* and the remote config is copied to the local backend

Moving Backends Summary

- ◆ To make this work, you need to use a two-step process:
 - Create the S3 bucket and DynamoDB table and deploy that code with a local backend
 - Add a remote backend configuration to it to use the S3 bucket and DynamoDB table
 - Run terraform init to copy your local state to S3
- ◆ To revert to a local state backend
 - Remove the backend configuration
 - Rerun terraform init to copy the Terraform state to the local disk
 - Run terraform destroy to delete the S3 bucket and DynamoDB table

Remote Backend Advantage

- ◆ A single S3 bucket and DynamoDB table can be shared across all your Terraform code
- ◆ You'll probably only need to do it once per AWS account
- ◆ After the S3 bucket exists, in the rest of your Terraform code, you can specify the backend configuration right from the start without any extra steps

Backend Limitation

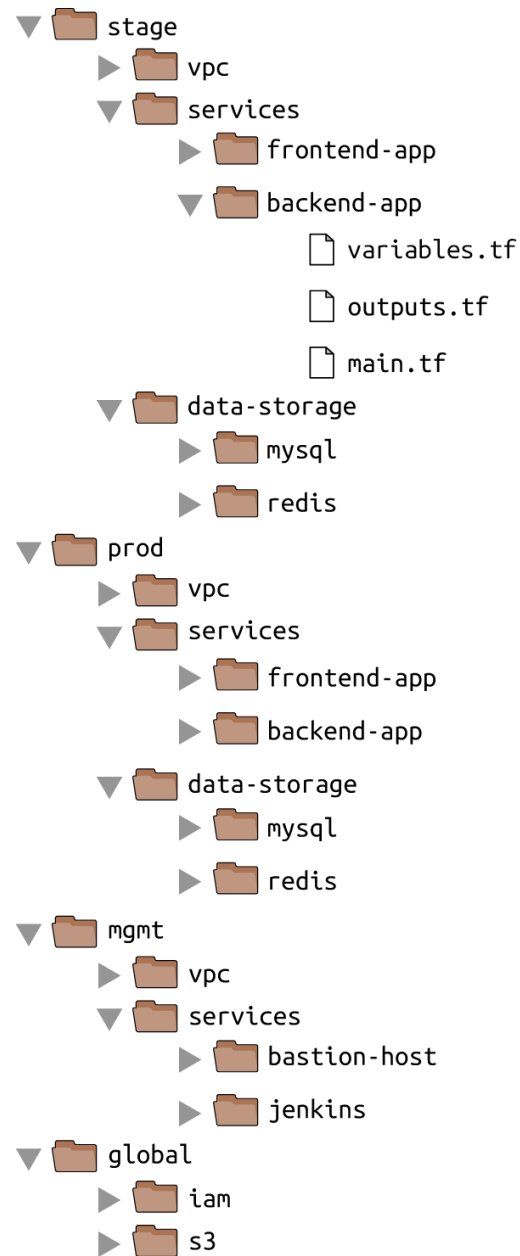
- ◆ Variables and references cannot be used in the `backend` block
- ◆ The following will **not** work

```
# This will NOT work. Variables aren't allowed in a backend configuration.
terraform {
  backend "s3" {
    bucket      = var.bucket
    region      = var.region
    dynamodb_table = var.dynamodb_table
    key         = "example/terraform.tfstate"
    encrypt     = true
  }
}
```

File Isolation

- ◆ Most secure approach is to have a folder for each configuration
 - Each folder can maintain its own version control for the *.tf files
 - Or a common repository can be used
 - Remember that the *.tf files are *source code*
- ◆ Each deployment has its own backend, local or remote.
 - This allows for isolation of all files
 - Allows for different access and authentication mechanisms
 - Eg. Different S3 buckets used as backends can have different policies

File Isolation Example



Workspaces Use Case

- ◆ If you already have a Terraform module deployed
 - you want to do some experiments with it
 - but you don't want your experiments to affect the state of the already deployed infrastructure
- ◆ Run *terraform workspace new* to deploy a new copy of the exact same infrastructure, but storing the state in a separate file

Workspace Specific Configurations

- ◆ You can even change how that module behaves based on the workspace you're in by reading the workspace name using the expression *terraform.workspace*

```
resource "aws_instance" "example" {  
  ami          = "ami-0c55b159cbfafa1f0"  
  instance_type = terraform.workspace == "default" ? "t2.medium" : "t2.micro"  
}
```

- ◆ Workspaces allow a fast and easy way to quickly spin up and tear down different versions of your code

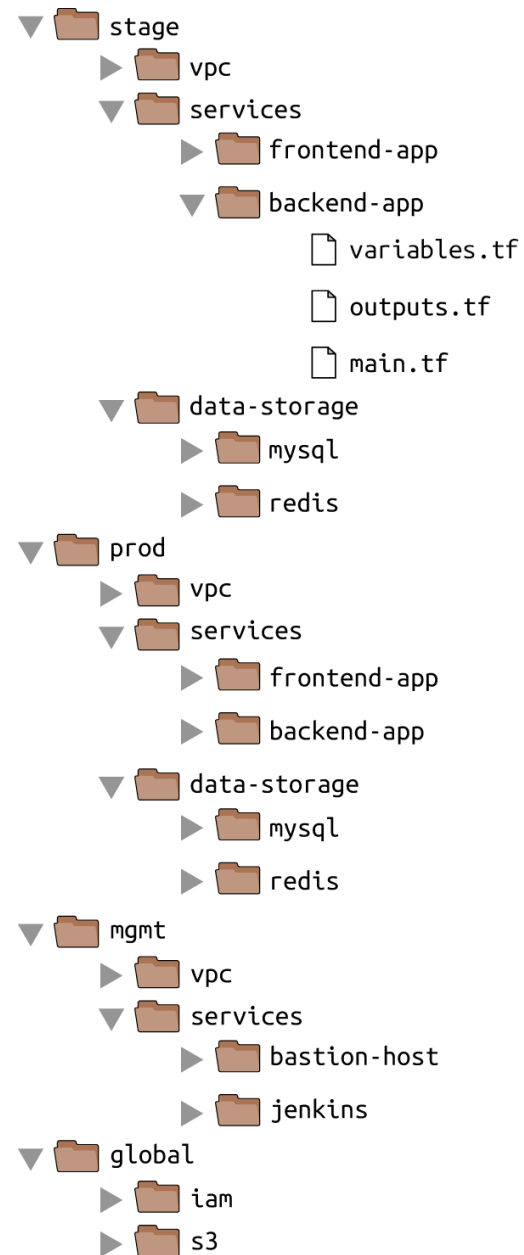
Workspace Drawbacks

- ◆ All workspace state files are stored in the same backend
 - They share same authentication and access controls which means they are not good for isolating
- ◆ Workspaces are not visible in the code or on the terminal unless you run terraform workspace commands
 - A module in one workspace looks exactly the same as a module deployed in 10 workspaces
 - This makes maintenance more difficult, because you don't have a good picture of your infrastructure
- ◆ Workspaces can be fairly error prone
 - The lack of visibility makes it easy to forget what workspace you're in and accidentally make changes in the wrong one

Isolation via File Layout

- ◆ To achieve full isolation between environments:
 - Put the Terraform configuration files for each environment into a separate folder
 - For example, all of the configurations for the staging environment can be in a folder called stage
 - All the configurations for the production environment can be in a folder called prod
- ◆ Configure a different backend for each environment, using different authentication mechanisms and access controls
 - Each environment could live in a separate AWS account with a separate S3 bucket as a backend

Typical Project File Layout



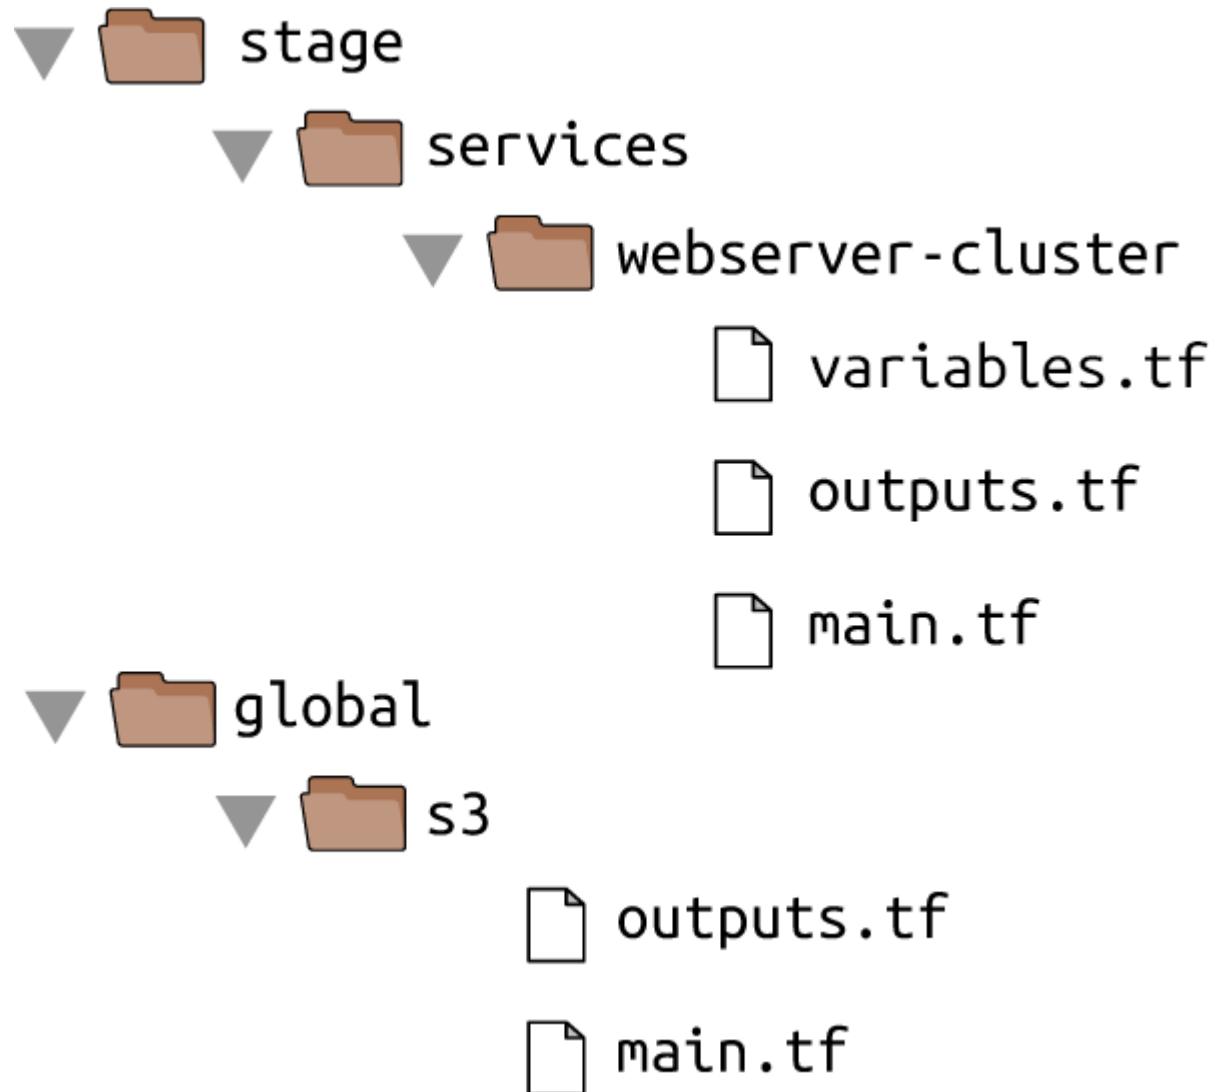
Isolation via File Layout

- ◆ At the top level, there are separate folders for each “environment”
 - **stage** : An environment for preproduction workloads (testing)
 - **prod** : An environment for production workloads (user facing apps)
 - **mgmt** : An environment for DevOps tooling (Jenkins etc.)
 - **global** : Resources that are used across all environments (S3, IAM)
- ◆ Within each environment, there are separate folders for each “component”:
 - **vpc** : Network topology for this environment
 - **services** : Apps or microservices to run in this environment - each app could have its own folder to isolate it
 - **data-storage** : The data stores to run in this environment, such as MySQL or Redis

Isolation via File Layout

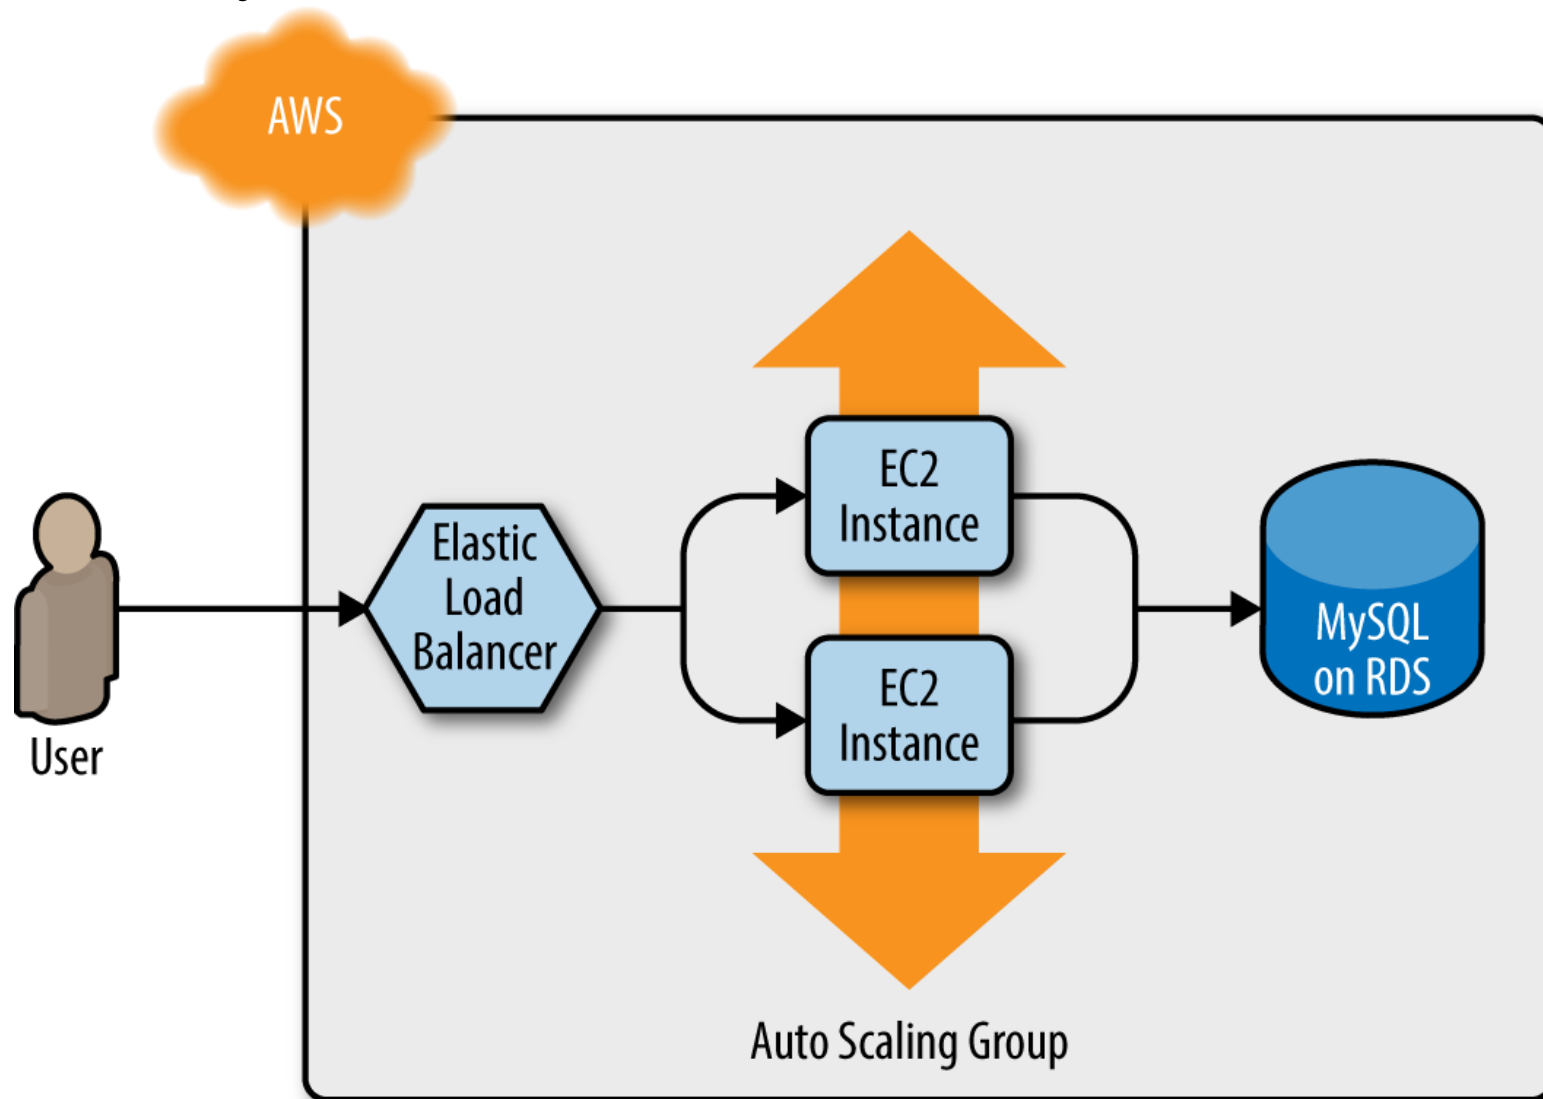
- ◆ Within each component are the actual Terraform configuration files with the following naming conventions:
 - **variables.tf** : Input variables
 - **outputs.tf** : Output variables
 - **main.tf** : The resources
- ◆ Terraform looks for files in the current directory with the .tf extension
 - Using a consistent, predictable naming convention makes code easier to browse
 - Then you always know where to look to find a variable, output, or resource. If individual Terraform files are

Rarranged Sample Code



The "terraform_remote_state" Data Source

- ◆ Assume that the web server cluster needs to communicate with a MySQL database

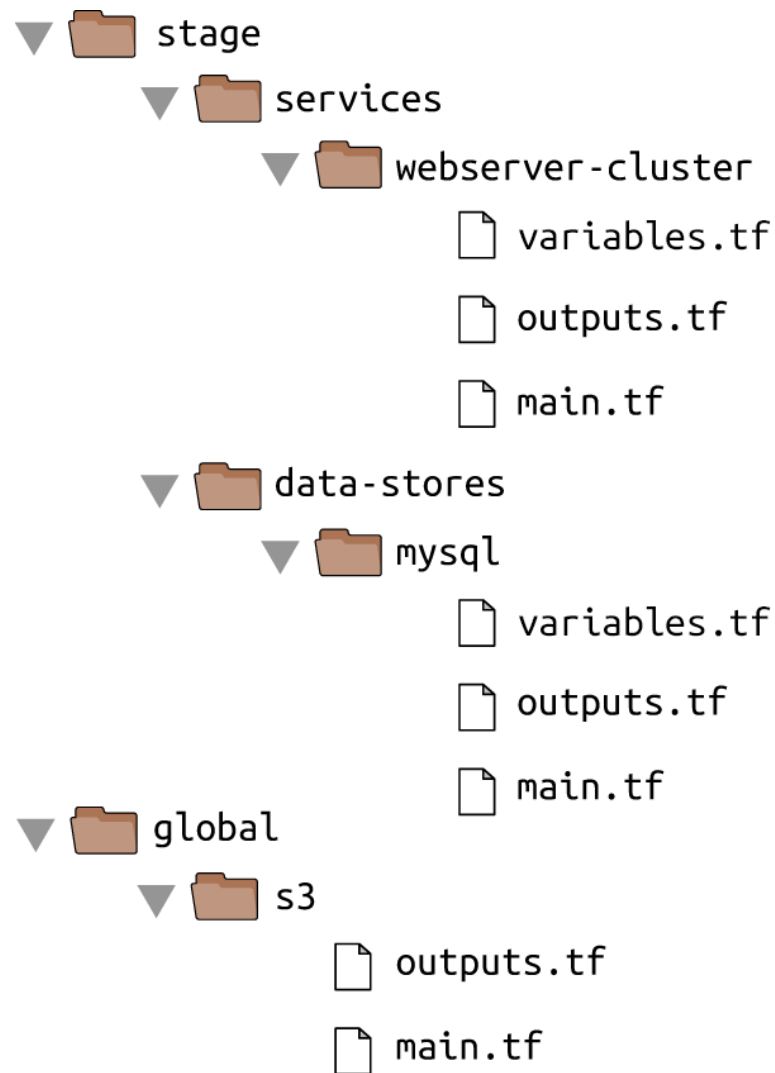


Deployment Consideration

- ◆ The MySQL database should probably be managed with a different set of configuration files as the web server cluster
 - Updates will probably be deployed to the web server cluster frequently
 - What to avoid accidentally breaking the database when doing an update

Deployment Consideration

- ◆ Isolate the MySQL configurations in a data-stores folder



Keeping Secrets I

- ◆ One of the parameters that you must pass to the `aws_db_instance` resource is the master password to use for the database
 - This should not be in the code in plain text
 - There are two other options
- ◆ Read the secret from a secret store - there are multiple secrets managers
 - AWS Secrets Manager and the `aws_secretsmanager_secret_version` data source (shown in the example code)
 - AWS Systems Manager Parameter Store and the `aws_ssm_parameter` data source
 - AWS Key Management Service (AWS KMS) and the `aws_kms_secrets` data source
 - Google Cloud KMS and the `google_kms_secret` data source
 - Azure Key Vault and the `azurerm_key_vault_secret` data source
 - HashiCorp Vault and the `vault_generic_secret` data source

Using AWS Secrets Manager

```
resource "aws_db_instance" "example" {  
  identifier_prefix = "terraform-up-and-running"  
  engine           = "mysql"  
  allocated_storage = 10  
  instance_class   = "db.t2.micro"  
  name             = "example_database"  
  username         = "admin"  
  
  password = data.aws_secretsmanager_secret_version.db_password.secret_string  
}  
  
data "aws_secretsmanager_secret_version" "db_password" {  
  secret_id = "mysql-master-password-stage"  
}
```

Keeping Secrets II

- ◆ Other option is to manage them completely outside of Terraform
 - Then pass the secret into Terraform via an environment variable.
 - In the code below, there is no default since it's a secret

```
variable "db_password" {  
  description = "The password for the database"  
  type        = string  
}  
  
export TF_VAR_db_password="(YOUR_DB_PASSWORD)"  
$ terraform apply
```

- ◆ A known weakness of Terraform:
 - The secret will be stored in the Terraform state file in plain text
 - The only solution is to lock down and encrypt the state files

Final Notes

- ◆ Correct isolation, locking and state must be a priority
 - Bugs in a program only break a part of an app
 - Bugs in infrastructure can have catastrophic effects and result in whole systems crashing and becoming unworkable
- ◆ Infrastructure has to be planned and incrementally tested
 - We never code infrastructure "on the fly"
- ◆ We never experiment with infrastructure in a production environment
 - Always work in a sandbox
 - With IaaS, this is easily done

Lab 3-3

- ◆ Please do Lab 3-3