

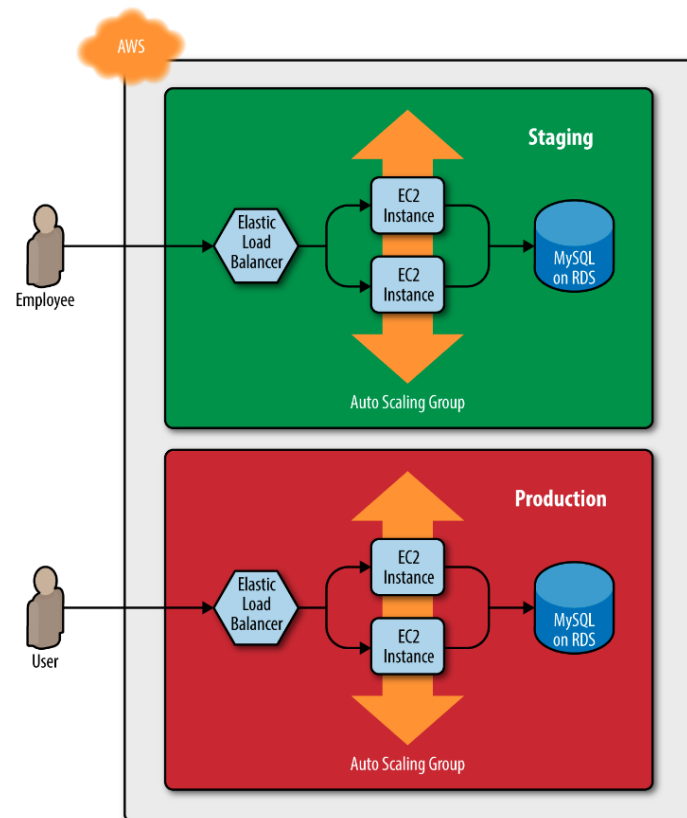
Reusable Infrastructure with Modules

The Plan

- ◆ Module basics
- ◆ Module inputs
- ◆ Module locals
- ◆ Module outputs
- ◆ Module gotchas
- ◆ Module versioning

Multiple Environments

- ◆ Cloud advantage: create multiple copies of the same environment
 - Production, Development, Test
 - Environments need to be similar if not identical
- ◆ We don't want to be able to re-use Terraform code across environments
 - DRY Principle: "Do not repeat yourself"

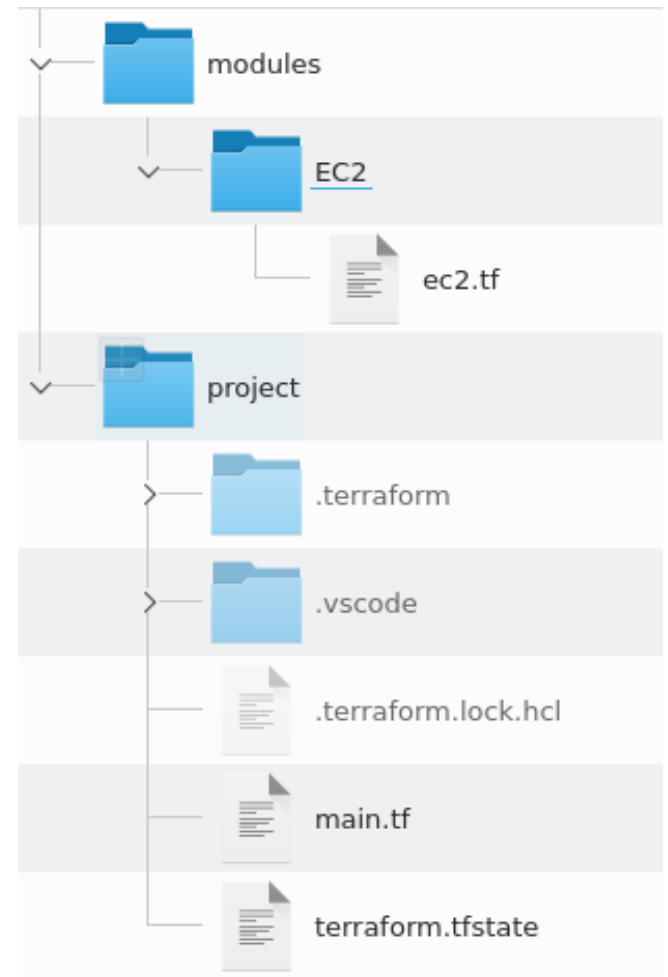


Module Basics

- ◆ Any folder containing terraform files is a module
 - There are no special declarations or syntax required
 - Modules are containers for multiple resources that are used together
 - Modules are the primary strategy used to package and reuse terraform resources
- ◆ Every terraform configuration has at least one module
 - It is referred to as the "root" module
 - It consists of the terraform files in the main working directory
- ◆ Modules (usually the root) may import other or "call" other modules
 - Modules that are being called are called "child" modules

Calling Modules Example

- ◆ If we are creating the same resource in multiple configurations, we can put it into a module
 - In this example, the demo project uses a module in the *modules/EC2* folder to create an EC2 instance to function like a webserver
 - The folder structure looks like this:



Calling Modules Example

- ◆ The EC2 Module code is familiar

```
resource "aws_instance" "alpha" {  
  ami = "ami-047a51fa27710816e"  
  instance_type = "t2.micro"  
  tags = {  
    source = "EC2 Module"  
  }  
}
```

- ◆ Calling it as a module is straightforward

```
module "VM-1" {  
  source = "../modules/EC2"  
}
```

- ◆ The problem is that this module is not easily reusable because the *ami* and *instance type* are hard-coded into the module code
 - We need to parameterize the module to make it reusable

Module Inputs

- ◆ Following the example of calling functions in a programming language, we want to be able to pass values to a module as parameters
- ◆ In the example, we want to parameterize the module code by adding three variables:

```
variable "ami_type" {  
    type = string  
    default = "ami-00399ec92321828f5"  
}  
  
variable "inst_type" [  
    type = string  
    default = "t2.nano"  
]  
  
variable VM_name {  
    type = string  
    default = "EC2-Module"  
}
```

Inside the EC2 Module

- ◆ We can parameterize the EC2 definitions

```
# In the EC2 Module

resource "aws_instance" "alpha" {
  ami = var.ami_type
  instance_type = var.inst_type
  tags = {
    source = "EC2 Module"
    Name = var.VM_name
  }
}
```


Passing Parameters

- ◆ We can pass the values to the variables when the EC2 module is called

```
module "VM-1" {  
    source = "../modules/EC2"  
    ami_type = "ami-00399ec92321828f5"  
    inst_type = "t2.micro"  
    VM_name = "alpha"  
}  
  
module "VM-2" {  
    source = "../modules/EC2"  
    ami_type = "ami-00399ec92321828f5"  
    inst_type = "t2.small"  
    VM_name = "beta"  
}
```

Variable Passthroughs

- ◆ We don't want to have to hardcode the parameters so we can use root module variables to pass the values through

```
variable machine_names {  
  type = list(string)  
  default = ["Singleton"]  
}  
  
variable machine_amis {  
  type = list(string)  
  default = ["ami-077e31c4939f6a2f3"]  
}  
  
variable machine_types {  
  type = list(string)  
  default = ["t2.nano"]  
}
```

```
module "VM-1" {  
  source = "../modules/EC2"  
  ami_type = var.machine_amis[0]  
  inst_type = var.machine_types[0]  
  VM_name = var.machine_names[0]  
}
```

Using "count" with Modules

- ◆ We can create multiple copies of modules with count

```
module "VM" {  
  count = length(var.machine_names)  
  source = "../modules/EC2"  
  ami_type = var.machine_amis[count.index]  
  inst_type = var.machine_types[count.index]  
  VM_name = var.machine_names[count.index]  
}
```

```
# terraform.tfvars  
|  
machine_names = ["alpha","beta"]  
machine_types = ["t2.small","t2.nano"]  
machine_amis = ["ami-077e31c4939f6a2f3","ami-00399ec92321828f5"]
```

Combining Modules

- ◆ Generally, we need to use more than one module because we need more than one type of resource
- ◆ In addition to the Ec2 instance module, we define a security group module

```
resource "aws_security_group" "app_port" {
  description = "Security group to allow access app instance"
  ingress {
    description = "OpenPort"
    from_port   = var.access_port
    to_port     = var.access_port
    protocol    = "tcp"
    cidr_blocks = ["0.0.0.0/0"]
  }

  egress {
    from_port   = 0
    to_port     = 0
    protocol    = "-1"
    cidr_blocks = ["0.0.0.0/0"]
  }

  tags = {
    Name = var.sg_name
  }
}
```

Combining Modules

- ◆ The security group has its own variables and called just like the EC2 module

```
variable "access_port" {  
    description = "Access port to use for the application"  
    type = number  
    default = 80  
}  
  
variable "sg_name" {  
    description = "The name of the security group"  
    type = string  
    default = "My SG"  
}
```

```
module "VM" {  
    count = length(var.machine_names)  
    source = "../modules/EC2"  
    ami_type = var.machine_amis[count.index]  
    inst_type = var.machine_types[count.index]  
    VM_name = var.machine_names[count.index]  
}  
  
module "SG" {  
    source = "../modules/SGroup"  
}
```

Module Return Value

- ◆ To use the security group created by the SG module, we need to be able to reference it
 - An output variable is passed back to the calling module with the id of the security group

```
output "secgps" {  
    value = aws_security_group.app_port.id  
    description = "Returns the id of the security group"  
}
```

- And an input variable for the EC2 module so it can be passed the security group reference

```
variable "sg_groups" {  
    description = "Associated security groups"  
    type = string  
}
```

- And add the reference to the resource definition

```
resource "aws_instance" "alpha" {  
    ami = var.ami_type  
    instance_type = var.inst_type  
    tags = {  
        source = "EC2 Module"  
        Name = var.VM_name  
    }  
    vpc_security_group_ids = [var.sg_groups]  
}
```

Module Return Value

- ◆ We can reference the value returned by the module with the following syntax
 - Reminder that the module name is the value we create when we *call* the module

```
1 module.< MODULE_NAME >.< OUTPUT_NAME >
```

```
module "VM" {  
    source = "../modules/EC2"  
    ami_type = var.machine_ami  
    inst_type = var.machine_type  
    VM_name = var.machine_name  
    sg_groups = module.SG.secgps  
}
```

Module Gotchas - Paths

- ◆ The hard-coded file paths are interpreted as relative to the current working directory
 - The problem is that this will not work if we are working with a module in a different directory
- ◆ To solve this issue, you can use an expression known as a path reference, which is of the form `path.<TYPE>`. Terraform supports the following types of path references:
 - `path.module` : Returns the file system path of the module where the expression is defined
 - `path.root` : Returns the file system path of the root module
 - `path.cwd` : Returns the file system path of the current working directory, usually the same as `path.root`

Module Path

- ◆ In this example, the template file is located with a path relative to the module, but if we hard-code the path, it will be interpreted as relative to the current working directory
 - By using the `path.module` construct, we ensure the file reference remains relative to the module

```
1 data "template_file" "user_data" {  
2   template = file("${path.module}/user-data.sh")  
3  
4   vars = {  
5     server_port = var.server_port  
6     db_address  = data.terraform_remote_state.db.outputs.address  
7     db_port     = data.terraform_remote_state.db.outputs.port  
8   }  
9 }
```

Module Gotcha - Inline Blocks

- ◆ The configuration for some Terraform resources can be defined either as inline blocks or as separate resources
 - When creating a module, you should always prefer using a separate resource
- ◆ Inline block example

```
1 resource "aws_security_group" "alb" {
2   name = "${var.cluster_name}-alb"
3
4   ingress {
5     from_port    = local.http_port
6     to_port      = local.http_port
7     protocol     = local.tcp_protocol
8     cidr_blocks  = local.all_ips
9   }
10
11   egress {
12     from_port    = local.any_port
13     to_port      = local.any_port
14     protocol     = local.any_protocol
15     cidr_blocks  = local.all_ips
16   }
17 }
```

Separate Resource

- ◆ You should change this module to define the exact same ingress and egress rules by using separate `aws_security_group_rule` resources

```
1 resource "aws_security_group" "alb" {
2   name = "${var.cluster_name}-alb"
3 }
4
5 resource "aws_security_group_rule" "allow_http_inbound" {
6   type           = "ingress"
7   security_group_id = aws_security_group.alb.id
8
9   from_port    = local.http_port
10  to_port      = local.http_port
11  protocol     = local.tcp_protocol
12  cidr_blocks  = local.all_ips
13 }
14
15 resource "aws_security_group_rule" "allow_all_outbound" {
16   type           = "egress"
17   security_group_id = aws_security_group.alb.id
18
19   from_port    = local.any_port
20   to_port      = local.any_port
21   protocol     = local.any_protocol
22   cidr_blocks  = local.all_ips
23 }
```

Inline Blocks

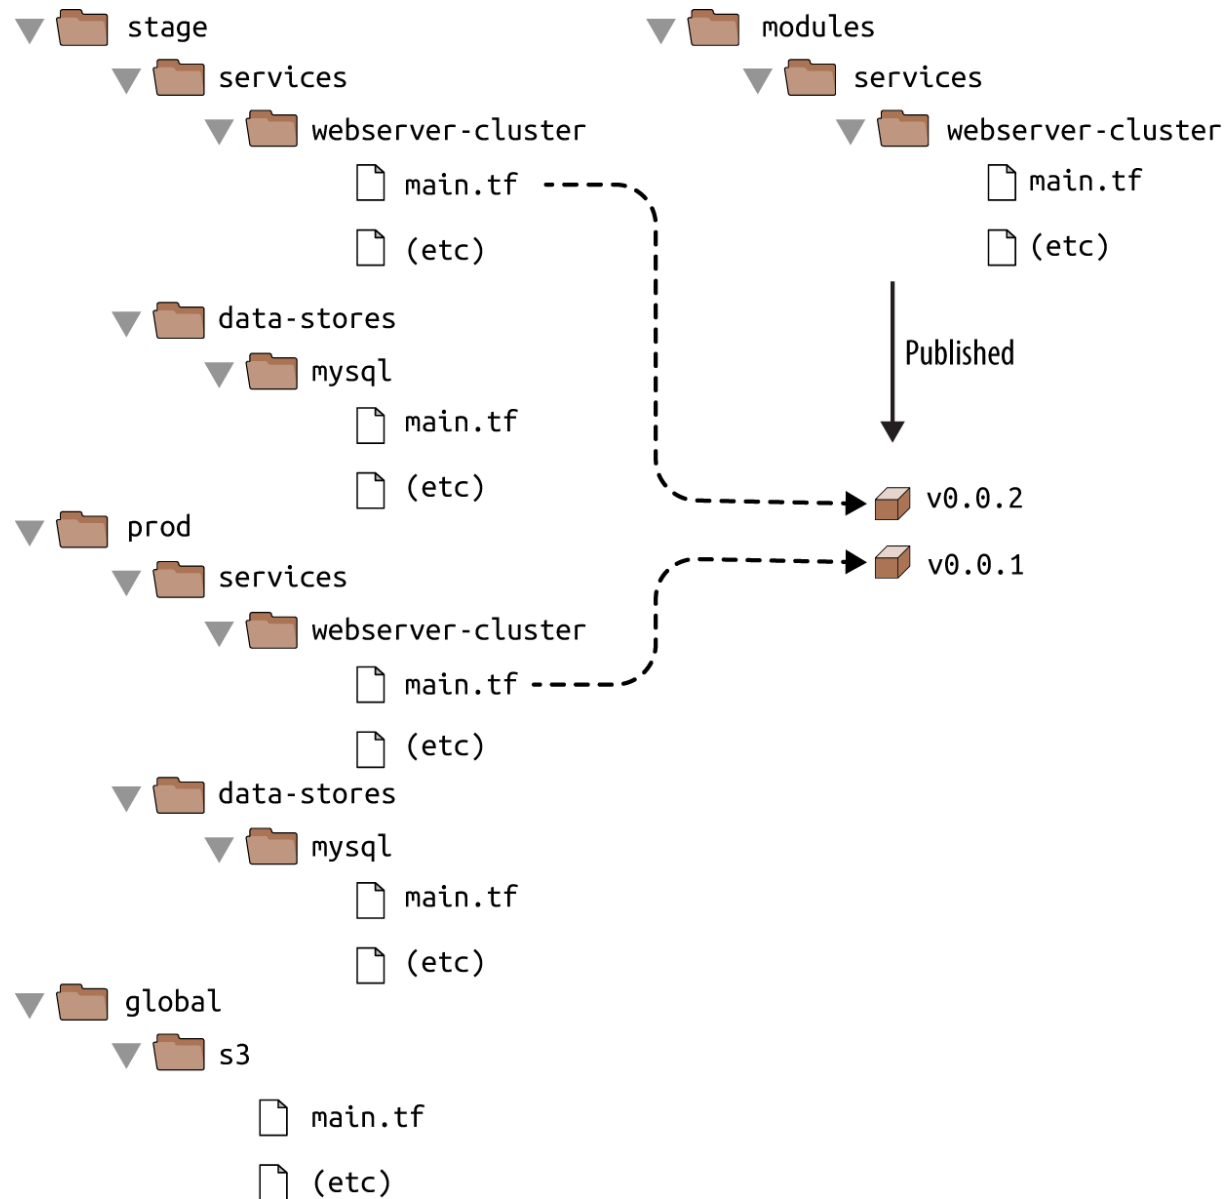
- ◆ Using a mix of inline blocks and separate resources may cause errors where routing rules conflict and overwrite one another
 - Use one or the other
 - When creating a module, you should always try to use a separate resource instead of the inline block
 - This allows for more flexible modules
- ◆ For example, changing a security group rule to allow a testing port is easier to do with a separate resource than having to edit inline blocks

```
1 resource "aws_security_group_rule" "allow_testing_inbound" {  
2     type          = "ingress"  
3     security_group_id = module.webserver_cluster.alb_security_group_id  
4  
5     from_port    = 12345  
6     to_port      = 12345  
7     protocol     = "tcp"  
8     cidr_blocks  = ["0.0.0.0/0"]  
9 }
```

Module Versioning

- ◆ If the staging and production environment point to the same module folder, any change in that folder will affect both environments on the very next deployment
 - This creates a coupling between environments and modules that can cause problems
- ◆ To solve this problem, we use a standard build management technique of using versions
 - As changes are made to a module, releases or versions of that module are published
 - Part of the configuration of any Terraform configuration plan is identification of which version of a module to include

Module Versioning Layout



Module Versioning

- ◆ An effective strategy is to use a repository tool like git and GitHub to publish releases of a module
 - Then the appropriate "release" of a module can be used

```
1 module "webserver_cluster" {
2   source = "github.com/foo/modules//webserver-cluster?ref=v0.0.1"
3
4   cluster_name      = "webservers-stage"
5   db_remote_state_bucket = "(YOUR_BUCKET_NAME)"
6   db_remote_state_key   = "stage/data-stores/mysql/terraform.tfstate"
7
8   instance_type = "t2.micro"
9   min_size      = 2
10  max_size      = 2
11 }
```

Semantic Versioning

- ◆ A common versioning scheme is "semantic versioning"
 - The format is MAJOR.MINOR.PATCH (e.g., 1.0.4)
 - There are specific rules on when you should increment each part of the version number
- ◆ MAJOR version increments when you make incompatible API changes
- ◆ MINOR version increments when you add functionality in a backward-compatible manner
- ◆ PATCH version increments when you make backward-compatible bug fixes