

Production-Grade Terraform Code

Production Infrastructure is Hard



What's Involved?

- ◆ Servers
- ◆ Data stores
- ◆ Load balancers
- ◆ Security
- ◆ Monitoring and alerting tools
- ◆ Building pipelines
- ◆ All the other pieces of your technology that are necessary to run a business

Goals

- ◆ Your infrastructure won't fall over if traffic goes up
- ◆ Not lose your data if there's an outage
- ◆ Not allow that data to be compromised when hackers try to break in
- ◆ If that is not achieved
 - Your company can go out of business.

Estimates

- ◆ Deploy a service fully managed by a third party
 - Such as running MySQL using the AWS Relational Database Service (RDS)
 - **One to two weeks** to get that service ready for production
- ◆ Your own stateless distributed app
 - Such as a cluster of Node.js apps
 - They store all their data in RDS
 - On top of an AWS Auto Scaling Group (ASG)
 - About **two to four weeks** to get ready for production

Bigger Estimates

- ◆ Your own stateful distributed app
 - Such as an Amazon Elasticsearch (Amazon ES) cluster
 - On top of an ASG and stores data on local disks
 - **Two to four months** to get ready for production
- ◆ build out your entire architecture including
 - All of your apps
 - Data stores
 - Load balancers
 - Monitoring
 - Alerting
 - Security
 - About **6 to 36 months** of work
 - Companies typically closer to six months
 - Larger companies typically taking several years

Estimate Summary

Type of infrastructure	Example	Time estimate
Managed service	Amazon RDS	1–2 weeks
Self-managed distributed system (stateless)	A cluster of Node.js apps	2–4 weeks
Self-managed distributed system (stateful)	Amazon ES	2–4 months
Entire architecture	Apps, data stores, load balancers, monitoring, etc.	6–36 months

Production Grade Infrastructure

- ◆ Why does it take so long to build production-grade infrastructure?
- ◆ The production-grade infrastructure checklist
- ◆ Production-grade infrastructure modules
 - Small modules
 - Composable modules
 - Testable modules
 - Releasable modules
 - Beyond Terraform modules

Why it Takes a Long Time

- ◆ Software estimates are notoriously inaccurate.
 - Even worse for DevOps
 - Unforeseen problems occur that disrupt the process
- ◆ The industry is in its infancy
 - Cloud computing, DevOps and tools like Terraform are all recent innovations
 - Still changing rapidly and innovating - not stable yet
 - Most people do not have a depth of experience with them
- ◆ The process is prone to disruption
 - The integration of all the details into a smooth process is still not stable
 - Diagnostic methods to solve problems are still maturing
 - There does not yet exist a body of experience and knowledge that covers a wide range of situations

The Complexity Problem

- ◆ Scaling up from simple environments where everything works easily to complex environments is always problematic
- ◆ Complexity is a major cause of software failure
- ◆ **Essential complexity** is when we are working with a complex problem - like building a whole corporate IT infrastructure
- ◆ **Accidental complexity** arises from the problems involved in using a specific set of tools and processes -- wrong tools and processes made the solution difficult
- ◆ **Environmental complexity** comes from disorganized work environments such as when there is a lack of processes or leadership
- ◆ Not being able to manage complexity effectively can often impact a project's timelines and costs as developers are overwhelmed by a flood of details

Infrastructure Checklist

- ◆ A major challenge is that different groups in a company have only a partial view of "going to production" based on their own focus of activity
- ◆ Infrastructure Checklist

Task	Description	Example tools
Install	Install the software binaries and all dependencies.	Bash, Chef, Ansible, Puppet
Configure	Configure the software at run-time. Includes port settings, TLS certs, service discovery, leaders, followers, replication, etc.	Bash, Chef, Ansible, Puppet
Provision	Provision the infrastructure. Includes servers, load balancers, network configuration, firewall settings, IAM permissions, etc.	Terraform, CloudFormation

Infrastructure Checklist

Deploy	Deploy the service on top of the infrastructure. Roll out updates with no downtime. Includes blue-green, rolling, and canary deployments.	Terraform, CloudFormation, Kubernetes, ECS
High availability	Withstand outages of individual processes, servers, services, data centers, and regions.	Multidatacenter, multiregion, replication, auto scaling, load balancing
Scalability	Scale up and down in response to load. Scale horizontally (more servers) and/or vertically (bigger servers).	Auto scaling, replication, sharding, caching, divide and conquer
Performance	Optimize CPU, memory, disk, network, and GPU usage. Includes query tuning, benchmarking, load testing, and profiling.	Dynatrace, valgrind, VisualVM, ab, Jmeter

Infrastructure Checklist

Networking	Configure static and dynamic IPs, ports, service discovery, firewalls, DNS, SSH access, and VPN access.	VPCs, firewalls, routers, DNS registrars, OpenVPN
Security	Encryption in transit (TLS) and on disk, authentication, authorization, secrets management, server hardening.	ACM, Let's Encrypt, KMS, Cognito, Vault, CIS
Metrics	Availability metrics, business metrics, app metrics, server metrics, events, observability, tracing, and alerting.	CloudWatch, DataDog, New Relic, Honeycomb
Logs	Rotate logs on disk. Aggregate log data to a central location.	CloudWatch Logs, ELK, Sumo Logic, Papertrail

Infrastructure Checklist

Backup and Restore	Make backups of DBs, caches, and other data on a scheduled basis. Replicate to separate region/account.	RDS, ElastiCache, replication
Cost optimization	Pick proper Instance types, use spot and reserved Instances, use auto scaling, and nuke unused resources.	Auto scaling, spot Instances, reserved Instances
Documentation	Document your code, architecture, and practices. Create playbooks to respond to incidents.	READMEs, wikis, Slack
Tests	Write automated tests for your infrastructure code. Run tests after every commit and nightly.	Terratest, inspec, server-spec, kitchen-terraform

Production Grade Modules

- ◆ Properties of production grade Terraform modules:
 - Small modules
 - Composable modules
 - Testable modules
 - Releasable modules
 - Beyond Terraform modules

Small Modules


- ◆ A well-known anti-pattern "the Kitchen Sink Module"
 - All code gets dumped into a single module
 - It just sort of "grows"
 - Often the root module
- ◆ Large modules (> several hundred lines of code) have downsides - they are:
 - **Slow** - the plan phase takes a long time to execute
 - **Insecure** - fine-grained permissions on resources becomes almost impossible
 - **Risky** - a single error can propagate across the infrastructure
 - **Difficult to understand** - they are a wall of (usually disorganized) text
 - **Difficult to review** - not just to read, but the plan output is overwhelming to read
 - **Difficult to test** - testing is hard enough already

General Design Principles

- ◆ In engineering, there are three basic design principles:
 - **Modularity** : systems are organized in self-contained packages or modules
 - **Cohesion** : each module should provide one service and only one module provides each service
 - **Coupling** : The connections between modules are only through defined interfaces
- ◆ These ideas should be implemented in Terraform production modules.

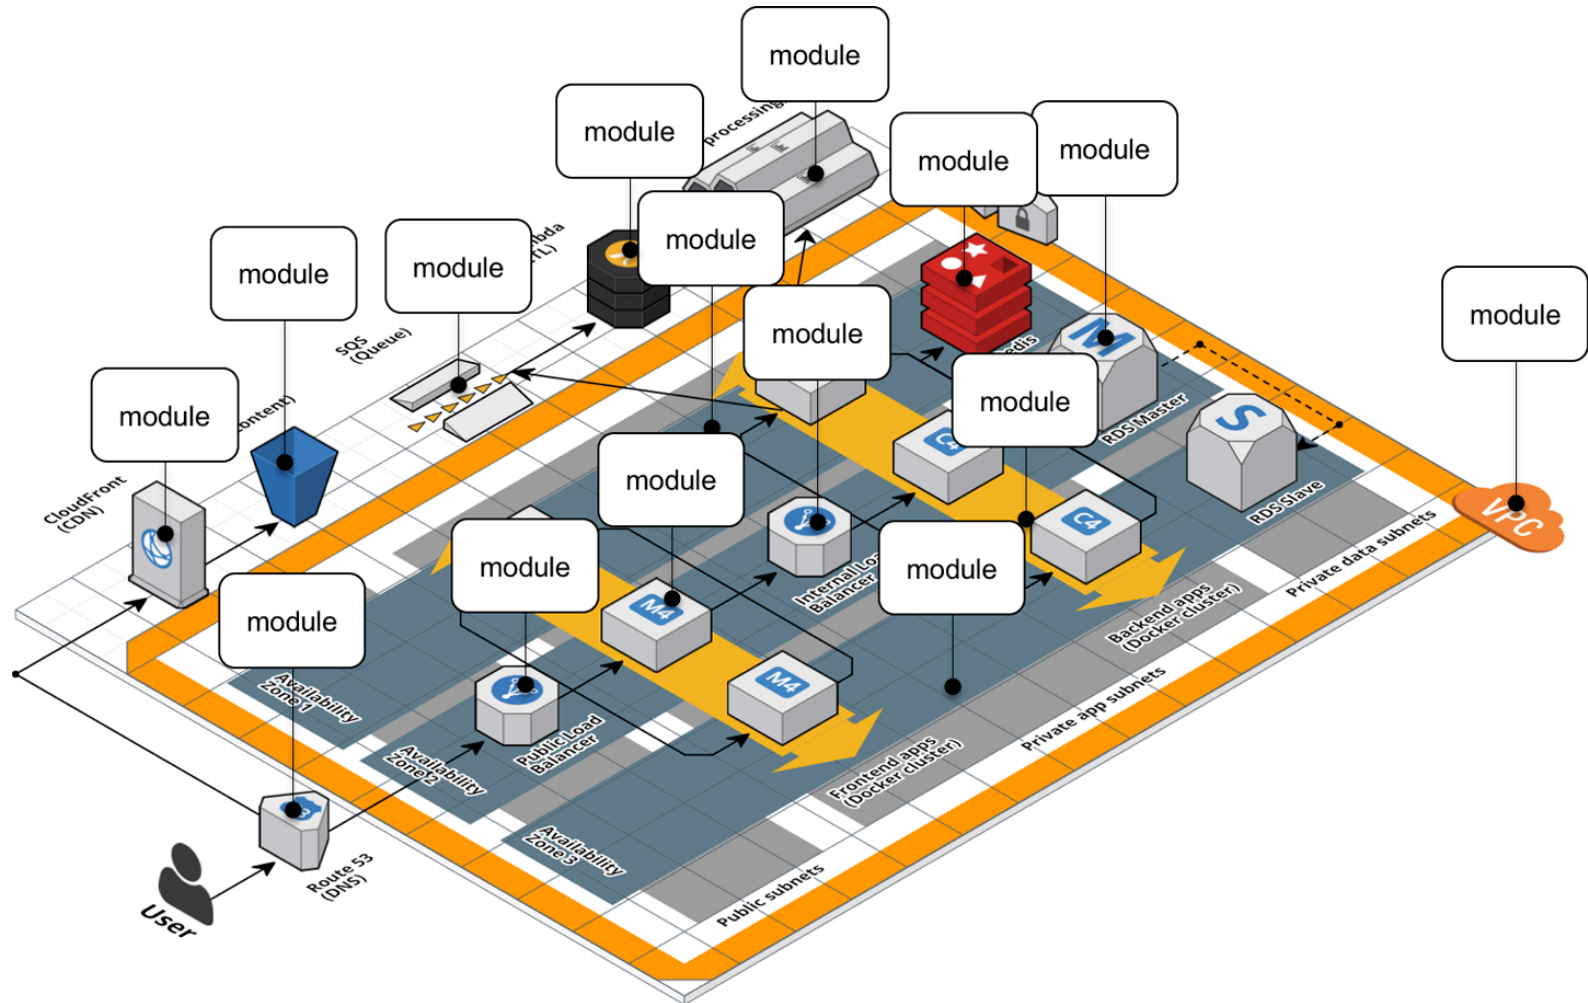
Architecture Example

- ◆ Given the following complex AWS architecture:



Architecture Example

- ◆ This suggests a module structure like this:



Refactoring Analysis

- ◆ The webserver cluster file manages three resources
 - Auto Scaling Group (ASG)
 - Application Load Balancer (ALB)
 - Hello, World app
- ◆ Refactoring is the process of changing the structure or organization of software without altering its functionality - like these smaller, cohesive modules
 - *modules/cluster/asg-rolling-deploy* - A generic, reusable, standalone module for deploying an ASG that can do a zero-downtime, rolling deployment
 - *modules/networking/alb* - A generic, reusable, standalone module for deploying an ALB
 - *modules/services/hello-world-app* - A module specifically for deploying the “Hello, World” app

Composable Modules

- ◆ Coupling requires we pass information between modules through interfaces
 - In Terraform, we do this by passing data through variables
- ◆ The core idea is to minimize side effects
 - Modules should avoid reading state information directly from the environment
 - Instead, state information is passed via input parameters
 - Essentially, there are "slots" in the module which are filled in with the information in the input variables
- ◆ Module that only produce a side effect and do not return a are difficult to use
 - If a module creates a resource, it should return a reference to the resource

Refactoring - Defining Variables

- ◆ In order to improve reuse for modules, input variables pass all the information the module needs to perform its function. For example, for the ASG Module:

```
variable "subnet_ids" {  
  description = "The subnet IDs to deploy to"  
  type        = list(string)  
}  
  
variable "target_group_arns" {  
  description = "The ARNs of ELB target groups in which to register Instances"  
  type        = list(string)  
  default     = []  
}  
  
variable "health_check_type" {  
  description = "The type of health check to perform. Must be one of: EC2, ELB."  
  type        = string  
  default     = "EC2"  
}  
  
variable "user_data" {  
  description = "The User Data script to run in each Instance at boot"  
  type        = string  
  default     = null  
}
```

Understanding the Variables

- ◆ *subnet_ids* : The original code had hard coded the default VPC and subnets, but by using this variable instead, the code can be used with any VPC or subnets
- ◆ *target_group_arns* : configures how the auto-scaling group integrates with load balancers
- ◆ *health_check_type* : also configures how the auto-scaling group integrates with load balancers
- ◆ Instead of the code just implementing as single fixed configuration, the variables allows you to use the ASG with a wide variety of use cases; e.g., no load balancer, one ALB, multiple NLBs, and so on

Modifying the Module Interface

- ◆ Once the variables are defined, the ASG module is modified to use the variables instead of any hardcoded resource references.

```
resource "aws_autoscaling_group" "example" {  
  # Explicitly depend on the launch configuration's name so each time  
  # it's replaced, this ASG is also replaced  
  name = "${var.cluster_name}-${aws_launch_configuration.example.name}"  
  
  launch_configuration = aws_launch_configuration.example.name  
  vpc_zone_identifier = var.subnet_ids  
  
  # Configure integrations with a load balancer  
  target_group_arns = var.target_group_arns  
  health_check_type = var.health_check_type  
  
  min_size = var.min_size  
  max_size = var.max_size  
  
  # Wait for at least this many instances to pass health checks before  
  # considering the ASG deployment complete  
  min_elb_capacity = var.min_size  
  
  # (...)  
}
```


The User Script

- ◆ The fourth variable, *user_data* , is for passing in a User Data script
- ◆ This allows us to deploy any application across an ALB

```
resource "aws_launch_configuration" "example" {  
  image_id      = var.ami  
  instance_type = var.instance_type  
  security_groups = [aws_security_group.instance.id]  
  user_data      = var.user_data  
  
  # Required when using a launch configuration with an auto scaling group.  
  # https://www.terraform.io/docs/providers/aws/r/launch\_configuration.html  
  lifecycle {  
    create_before_destroy = true  
  }  
}
```

Output Variables

- ◆ You'll also want to add a couple of useful output variables to `modules/cluster/asg-rolling-deploy/outputs.tf`
- ◆ Outputting this data makes the `asg-rolling-deploy` module even more reusable
 - Users of the module can add new behaviors, such as attaching custom rules to the security group

```
output "asg_name" {  
  value      = aws_autoscaling_group.example.name  
  description = "The name of the Auto Scaling Group"  
}  
  
output "instance_security_group_id" {  
  value      = aws_security_group.instance.id  
  description = "The ID of the EC2 Instance Security Group"  
}
```

Output Variables

- ◆ For similar reasons, several output variables can be added to *modules/networking/alb/outputs.tf* :
- ◆ We have defined return values that provide the name, arn and security group id that were created when the resources were instantiated

```
output "alb_dns_name" {
  value      = aws_lb.example.dns_name
  description = "The domain name of the load balancer"
}

output "alb_http_listener_arn" {
  value      = aws_lb_listener.http.arn
  description = "The ARN of the HTTP listener"
}

output "alb_security_group_id" {
  value      = aws_security_group.alb.id
  description = "The ALB Security Group ID"
}
```

Creating the "Hello World" App

- ◆ The last step is to convert the *webserver-cluster* module into a *hello-world-app* module that can deploy a “Hello, World” app using the *asg-rolling-deploy* and *alb* modules
- ◆ The resources left in *module/services/hello-world-app/main.tf* are:
 - *template_file* (for User Data)
 - *aws_lb_target_group*
 - *aws_lb_listener_rule*
 - *terraform_remote_state* (for the DB)
 - *aws_vpc*
 - *aws_subnet_ids*
- ◆ Add the following variable to *modules/services/hello-world-app/variables.tf*:

```
variable "environment" {  
  description = "The name of the environment we're deploying to"  
  type        = string  
}
```

Adding the "asg-rolling-deploy" Module

- ◆ Now, add the asg-rolling-deploy module that you created earlier to the hello-world-app module to deploy an ASG:

```
module "asg" {  
  source = "../../cluster/asg-rolling-deploy"  
  
  cluster_name = "hello-world-${var.environment}"  
  ami          = var.ami  
  user_data    = data.template_file.user_data.rendered  
  instance_type = var.instance_type  
  
  min_size      = var.min_size  
  max_size      = var.max_size  
  enable_autoscaling = var.enable_autoscaling  
  
  subnet_ids      = data.aws_subnet_ids.default.ids  
  target_group_arns = [aws_lb_target_group.asg.arn]  
  health_check_type = "ELB"  
  
  custom_tags = var.custom_tags  
}
```

Adding the "alb" Module

- ◆ Use of the input variable environment to enforce a naming convention, so all of your resources will be namespaced based on the environment (e.g., hello-world-stage, hello-world-prod)

```
module "alb" {  
  source = "../../networking/alb"  
  
  alb_name    = "hello-world-${var.environment}"  
  subnet_ids = data.aws_subnet_ids.default.ids  
}
```

Configure the ALB Target Group

- ◆ Configure the ALB target group and listener rule for this app
- ◆ Update the `aws_lb_target_group` resource in *modules/services/hello-world-app/main.tf* to use environment in its name:

```
resource "aws_lb_target_group" "asg" {
  name      = "hello-world-${var.environment}"
  port      = var.server_port
  protocol  = "HTTP"
  vpc_id    = data.aws_vpc.default.id

  health_check {
    path            = "/"
    protocol        = "HTTP"
    matcher         = "200"
    interval        = 15
    timeout         = 3
    healthy_threshold = 2
    unhealthy_threshold = 2
  }
}
```

Update Listener

- ◆ Update the `listener_arn` parameter of the `aws_lb_listener_rule` resource to point at the `alb_http_listener_arn` output of the ALB module:

```
resource "aws_lb_listener_rule" "asg" {
  listener_arn = module.alb.alb_http_listener_arn
  priority     = 100

  condition {
    path_pattern {
      values = ["*"]
    }
  }

  action {
    type              = "forward"
    target_group_arn = aws_lb_target_group.asg.arn
  }
}
```


Pass Through Values

- ◆ Pass through the important outputs from the *asg-rolling-deploy* and *alb* modules as outputs of the *hello-world-app* module:

```
output "alb_dns_name" {
  value      = module.alb.alb_dns_name
  description = "The domain name of the load balancer"
}

output "asg_name" {
  value      = module.asg.asg_name
  description = "The name of the Auto Scaling Group"
}

output "instance_security_group_id" {
  value      = module.asg.instance_security_group_id
  description = "The ID of the EC2 Instance Security Group"
}
```

Module Composition

- ◆ Composition is building up more complicated behavior for the “Hello, World” app from simpler parts (ASG and ALB modules)
- ◆ A fairly common pattern in terraform is that a configuration will have at least two types of modules:
- ◆ **Generic modules** : the basic building blocks of terraform code, reusable across a wide variety of use cases
- ◆ **Use-case-specific modules** : Combines multiple generic modules with some specific "glue" code to serve one specific use case such as deploying the “Hello, World” app

Testable Modules

- ◆ There is a lot of code in three modules:
 - *asg-rolling-deploy*
 - *alb*
 - *hello-world-app*
- ◆ The next step is to check that your code actually works.
 - These are not root modules, so we have to provide the infrastructure needed to run them
 - For example: backend configuration, provider, etc
- ◆ To do this, we use an "examples" folder that provides examples on how to use the module

Sample Example

- ◆ The code below is an example file to deploy an ASG of size 1

```
module "asg" {  
  source = "../../modules/cluster/asg-rolling-deploy"  
  
  cluster_name = var.cluster_name  
  ami          = "ami-0c55b159cbfafelf0"  
  instance_type = "t2.micro"  
  
  min_size      = 1  
  max_size      = 1  
  enable_autoscaling = false  
  
  subnet_ids = data.aws_subnet_ids.default.ids  
}  
  
data "aws_vpc" "default" {  
  default = true  
}  
  
data "aws_subnet_ids" "default" {  
  vpc_id = data.aws_vpc.default.id  
}
```

Test Assets

- ◆ The code on the previous slide provides:
- ◆ A manual test harness: you can use this example code to repeatedly deploy and undeploy manually to check that it works as you expect
- ◆ An automated test harness: The example code is also creates a framework for using automated testing tools
- ◆ Executable documentation: If you commit an example (including README.md) into version control then:
 - Other team members can use it to understand how the module works
 - They can try out the module without writing a line of code.
- ◆ Every module should have a corresponding example in the examples folder
 - There may be multiple examples showing different configurations and permutations of how that module can be used

Folder Structure

```
1 modules
2   L examples
3     L alb
4     L asg-rolling-deploy
5       L one-instance
6       L auto-scaling
7       L with-load-balancer
8       L custom-tags
9     L hello-world-app
10    L mysql
11  L modules
12    L alb
13    L asg-rolling-deploy
14    L hello-world-app
15    L mysql
16  L test
17    L alb
18    L asg-rolling-deploy
19    L hello-world-app
20    L mysql
```

Version Pinning

- ◆ You should pin all of your Terraform modules to a specific version of terraform using the `required_version` argument
- ◆ You want to avoid breaking a configuration because of a change in terraform itself
- ◆ At a bare minimum, you should require a specific major version of terraform:
- ◆ For production-grade code, it is recommended pinning the version even more strictly:

```
terraform {  
  # Require any 0.12.x version of Terraform  
  required_version = ">= 0.12, < 0.13"  
}  
  
terraform {  
  # Require Terraform at exactly version 0.12.0  
  required_version = "= 0.12.0"  
}
```

Releasable Modules

- ◆ Modules should be made available for use only after they have been "released"
 - This is what we saw in the Modules section
- ◆ Another option for releasing modules is to publish them in the Terraform Registry
 - The Public Terraform Registry resides at <https://registry.terraform.io/>
 - It includes hundreds of reusable, community-maintained, open source modules for AWS, Google Cloud, Azure, and many other providers

Publishing Requirements

- ◆ There are a few requirements to publish a module to the Public Terraform Registry
 - The module must live in a public GitHub repo
 - The repo must be named terraform-< PROVIDER >-< NAME >, where PROVIDER is the provider the module is targeting (e.g., aws) and NAME is the name of the module
 - The module must follow a specific file structure, including defining Terraform code in the root of the repo, providing a README.md, and using the convention of main.tf, variables.tf, and outputs.tf as filenames
 - The repo must use Git tags with semantic versioning (x.y.z) for releases

Using the Registry

- ◆ Terraform even supports a special syntax for consuming modules from the Terraform Registry
 - You can use a special shorter registry URL in the source argument and specify the version via a separate version argument using the following syntax:

```
module "<NAME>" {  
    source = "<OWNER>/<REPO>/<PROVIDER>"  
    version = "<VERSION>"  
}  
  
# for example  
  
module "vault" {  
    source = "hashicorp/vault/aws"  
    version = "0.12.2"  
}
```

Beyond Modules

- ◆ Sometimes non Terraform code has to be run from a terraform module or integrate with other tools or systems
- ◆ Sometimes we have to work around a limitation of terraform, like trying implement complicated logic
- ◆ Some Terraform "escape hatches" are:
 - Provisioners
 - Provisioners with *null_resource*
 - External Data Source

Provisioners

- ◆ Provisioners are used to execute scripts either on the local machine or a remote machine, typically to do the work of bootstrapping, configuration management, or cleanup
- ◆ There are several types of provisioners
 - *local-exec* : execute a script on the local machine
 - *remote-exec* : execute a script on a remote resource
 - *chef* : run Chef client on a remote resource
 - *file* : copy files to a remote resource

Provisioner Block

- ◆ Provisioners are added using the `provisioner` block

```
resource "aws_instance" "example" {  
  ami           = "ami-0c55b159cbfaffe1f0"  
  instance_type = "t2.micro"  
  
  provisioner "local-exec" {  
    command = "echo \"Hello, World from $(uname -smp)\""  
  }  
}
```

- ◆ Running `terraform apply` produces:

```
aws_instance.example: Creating...  
aws_instance.example: Still creating... [10s elapsed]  
aws_instance.example: Still creating... [20s elapsed]  
aws_instance.example: Provisioning with 'local-exec'...  
aws_instance.example (local-exec): Executing: ["/bin/sh" "-c" "echo \"Hello, World from $(unam  
e -smp)\""]  
aws_instance.example (local-exec): Hello, World from Linux x86_64 x86_64  
aws_instance.example: Creation complete after 23s [id=i-0da14be7bc14e434a]
```

Remote Provisioning

- ◆ Assume we want to provision an EC2 instance, we have to connect to the instance and authenticate to it
 - In this example we will use SSH
- ◆ We need a security group to allow SSH access:

```
resource "aws_security_group" "instance" {  
  ingress {  
    from_port = 22  
    to_port   = 22  
    protocol  = "tcp"  
  
    # To make this example easy to try out, we allow all SSH connections.  
    # In real world usage, you should lock this down to solely trusted IPs.  
    cidr_blocks = ["0.0.0.0/0"]  
  }  
}
```

Take a New Screenshot

Generate SSH Keys

```
# To make this example easy to try out, we generate a private key in Terraform.  
# In real-world usage, you should manage SSH keys outside of Terraform.  
resource "tls_private_key" "example" {  
  algorithm = "RSA"  
  rsa_bits  = 4096  
}  
  
resource "aws_key_pair" "generated_key" {  
  public_key = tls_private_key.example.public_key_openssh  
}
```

Creating and Connecting to the Instance

- ◆ The inline argument to pass a list of commands to execute, instead of a single command argument
- ◆ But we also have to configure the EC2 instance to use ssh

```
resource "aws_instance" "example" {  
  ami                = "ami-0c55b159cbfafa1f0"  
  instance_type      = "t2.micro"  
  vpc_security_group_ids = [aws_security_group.instance.id]  
  key_name           = aws_key_pair.generated_key.key_name  
  
  provisioner "remote-exec" {  
    inline = ["echo \"Hello, World from $(uname -smp)\""]  
  }  
  connection {  
    type      = "ssh"  
    host      = self.public_ip  
    user      = "ubuntu"  
    private_key = tls_private_key.example.private_key_pem  
  }  
}
```


Output

- ◆ Running *terraform apply* produces:

```
aws_instance.example (remote-exec): Connecting to remote host via SSH...
aws_instance.example (remote-exec): Host: 18.188.177.1
aws_instance.example (remote-exec): User: ubuntu
aws_instance.example (remote-exec): Password: false
aws_instance.example (remote-exec): Private key: true
aws_instance.example (remote-exec): Certificate: false
aws_instance.example (remote-exec): SSH Agent: true
aws_instance.example (remote-exec): Checking Host Key: false
aws_instance.example (remote-exec): Target Platform: unix
aws_instance.example (remote-exec): Connected!
aws_instance.example (remote-exec): Hello, World from Linux x86_64 x86_64
aws_instance.example: Creation complete after 43s [id=i-05b221a4369c4824a]
```

Data Scripts versus Provisioners

- ◆ Advantages to using a provisioner
 - Data scripts are limited to a length of 16KB, while provisioner scripts can be arbitrarily long
 - The Chef and other provisioners install, configure, and run on clients, which makes it easier to use configuration management tools
- ◆ The advantages to User Data scripts are:
 - You can use User Data scripts with ASGs, but Provisioners take effect only while Terraform is running and don't work with ASGs at all
 - The User Data script can be seen in the EC2 Console and you can find its execution log on the EC2 Instance itself, both of which are useful for debugging, neither of which is available with provisioners

Provisioners with "null_resource"

- ◆ Sometimes, you want to execute a provisioner without tying it to a specific resource
- ◆ we can use a *null_resource* which acts just like a normal terraform resource, except that it does not create anything
- ◆ The *null_resource* has an argument called *triggers*, which takes in a map of keys and values
- ◆ Whenever the values change, the *null_resource* will be recreated
- ◆ This forces any provisioners within it to be re-executed
- ◆ For example, the *uuid()* built-in function, which returns a new, randomly generated UUID each time it's called, within the *triggers* argument

```
resource "null_resource" "example" {  
  # Use UUID to force this null_resource to be recreated on every  
  # call to 'terraform apply'  
  triggers = {  
    uuid = uuid()  
  }  
  
  provisioner "local-exec" {  
    command = "echo \"Hello, World from $(uname -smp)\""  
  }  
}
```

Output from

- ◆ Every time you call terraform apply, the local-exec provisioner will execute:
- ◆ The output from the previous example is:

```
null_resource.example: Creating...  
null_resource.example: Provisioning with 'local-exec'...  
null_resource.example (local-exec): Executing: ["/bin/sh" "-c" "echo \"Hello, World from $(uname -smp)\""]  
null_resource.example (local-exec): Hello, World from Linux x86_64 x86_64  
null_resource.example: Creation complete after 0s [id=8600110595734706940]
```

External Data Source

- ◆ For situations where we want to execute a script to fetch some data and make that data available within the Terraform code itself
 - External data source allows an external command that implements a specific protocol to act as a data source
- ◆ The protocol is:
 - Data passes terraform to the external program using the query argument of the external data source
 - The external program can read in these arguments as JSON from stdin
 - The external program can pass data back to Terraform by writing JSON to stdout
 - Terraform code can then pull data out of this JSON by using the result output attribute of the external data source

External Example

- ◆ This example uses the external data source to execute a Bash script that echoes back to stdout any data it receives on stdin

```
data "external" "echo" {
  program = ["bash", "-c", "cat /dev/stdin"]
  query = {
    foo = "bar"
  }
}

output "echo" {
  value = data.external.echo.result
}

output "echo_foo" {
  value = data.external.echo.result.foo
}
```